

Twizzler : a Data-Centric OS for Non-Volatile Memory

Bittman Daniel, Alvaro Peter, Pankaj Mehra, Darrell D. E. Long, & Ethan L. Miller

In 2020 USENIX Annual Technical Conference

2020. 08. 10

Presentation by Han, Yejin

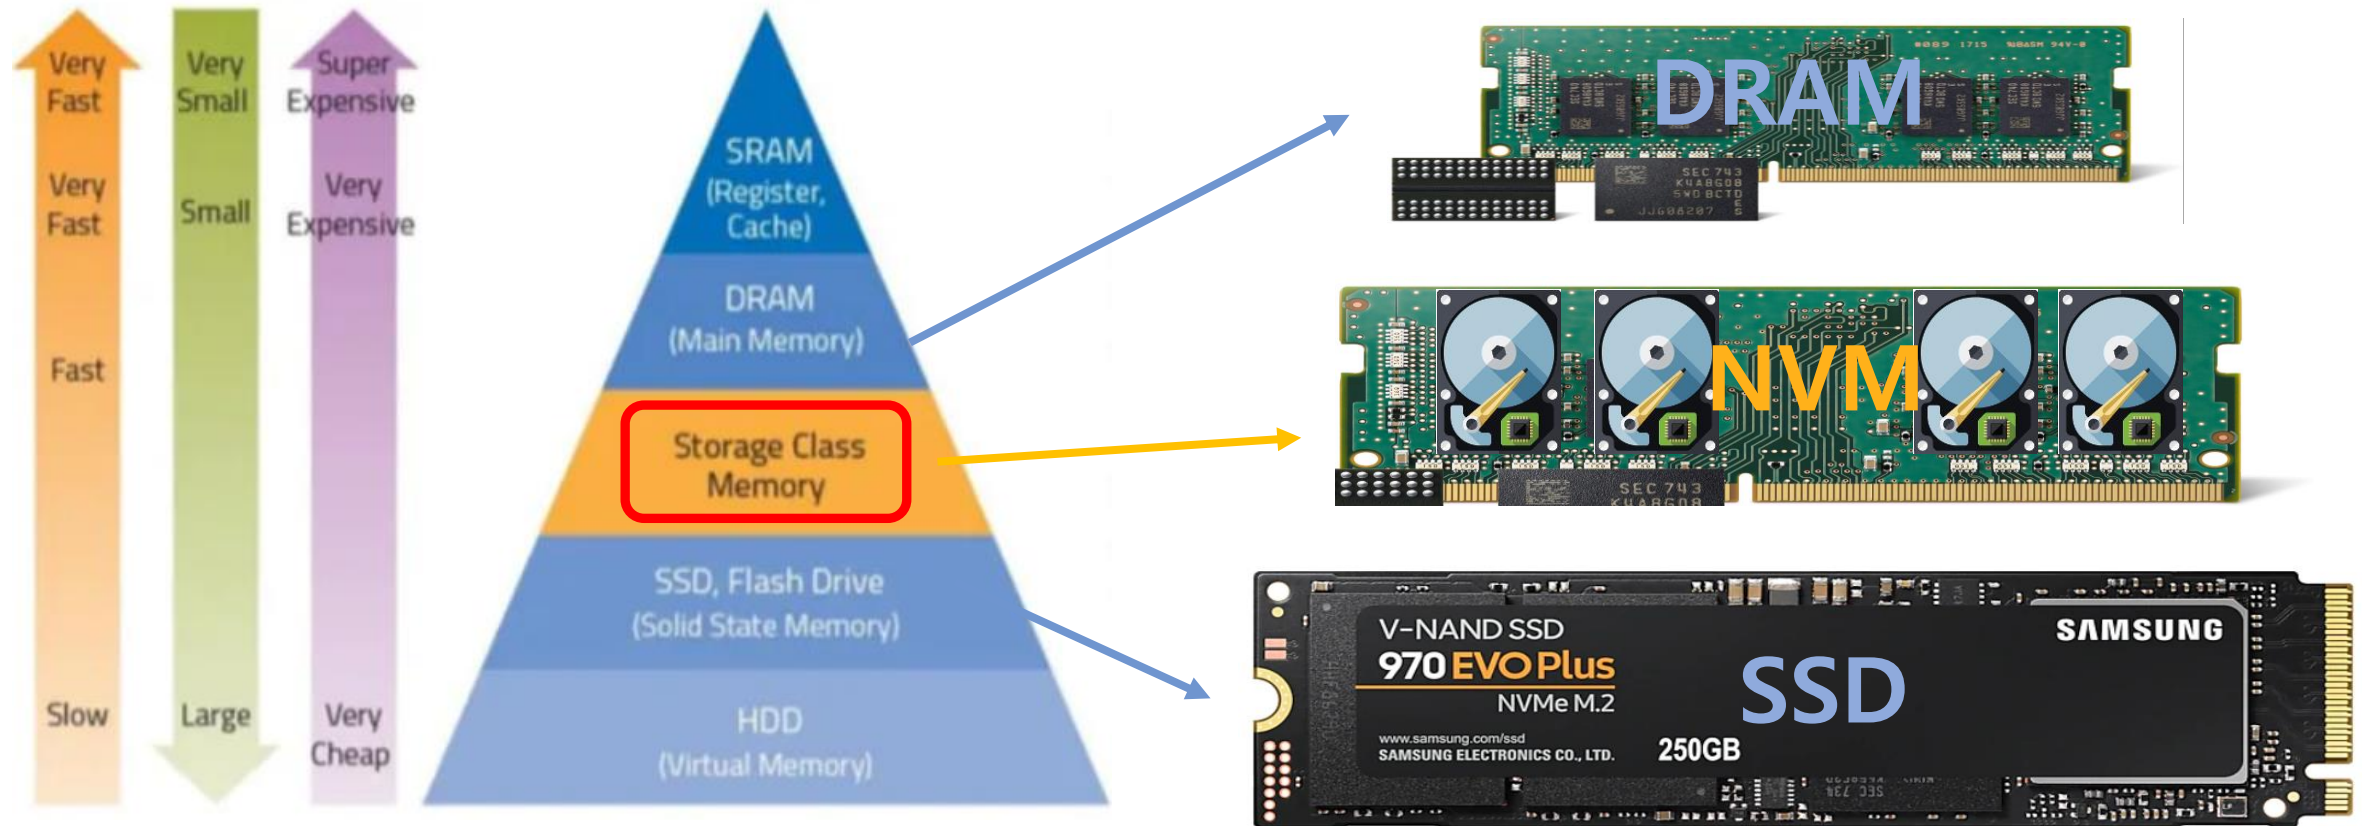
32164881@dankook.ac.kr

Contents

1. Introduction
2. Motivation
3. Twizzler
4. Evaluation
5. Conclusion

Hardware Trends

- Non-Volatile Memory
- Performance, Capacity, Cost



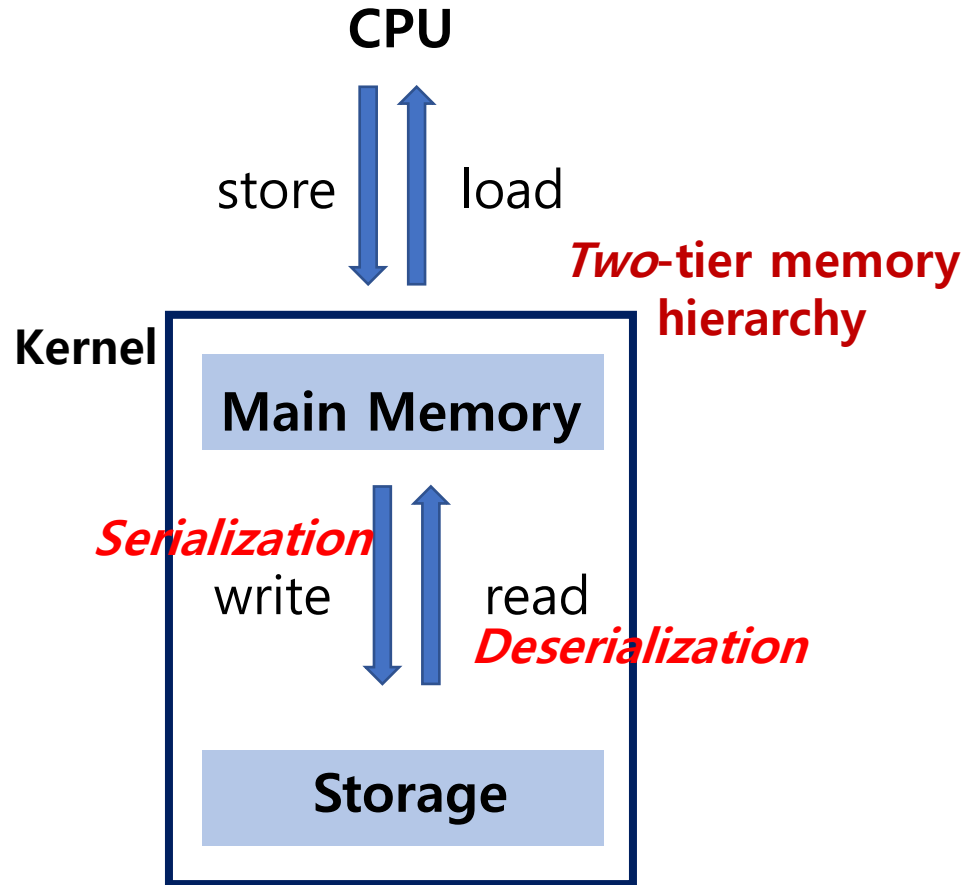
NVM characteristics

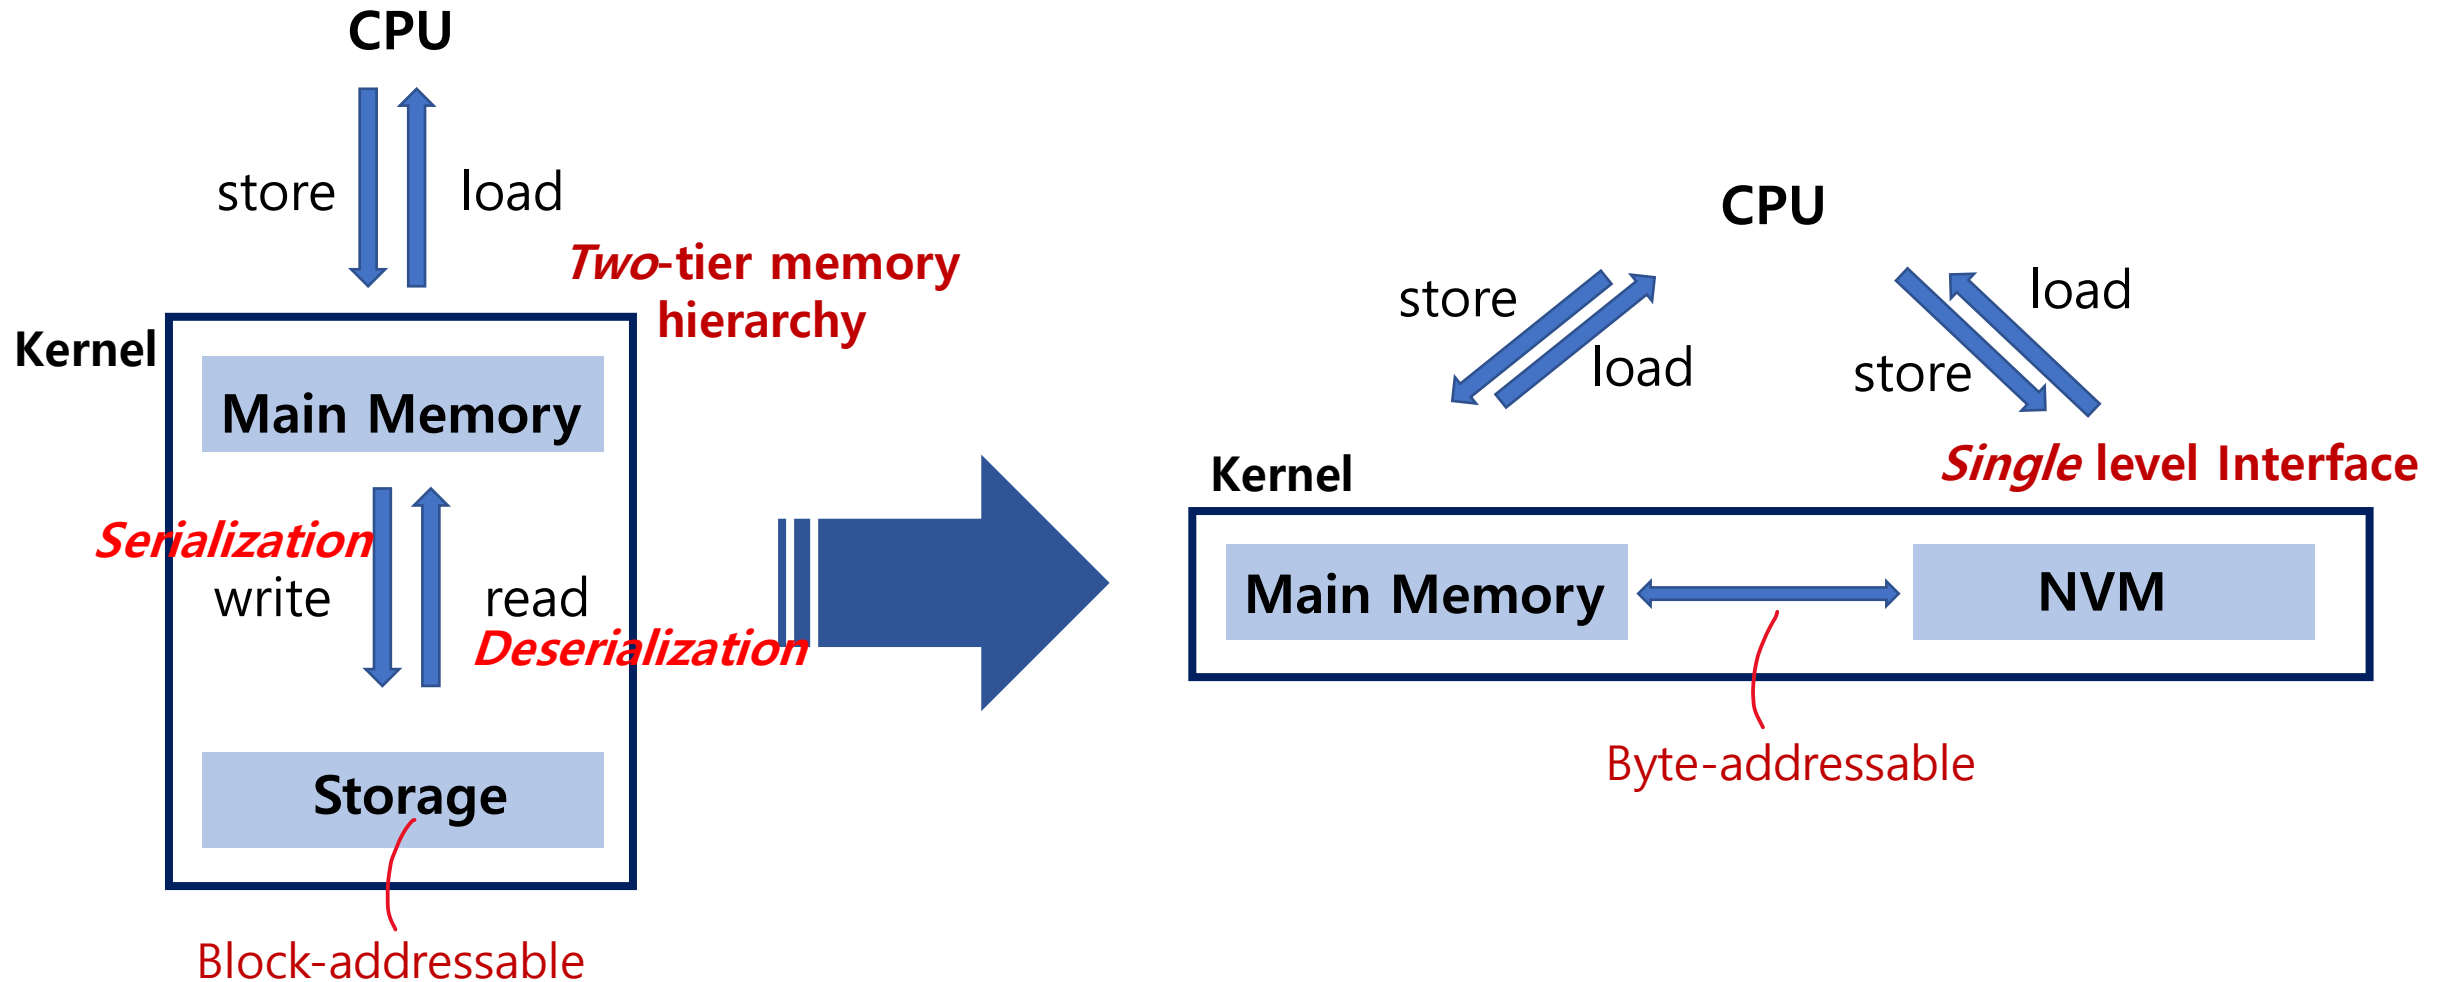
- ✓ Non-volatile
- ✓ R/W Performance
- ✓ Byte-addressable

COMPARISON OF HIGH-DENSITY MEMORY TECHNOLOGIES			
Attributes	DRAM	PCM	NAND
Non-Volatile	No	Yes	Yes
Erase Required	Bit	Bit	Block
Software	Simple	Simple	Complex
Power	-W/GB	100 → 500mW/die	-100mW/die
Write Bandwidth	-GB/s	1 → 100+ MB/s/die	10 → 100 MB/s/die
Write Latency	-20-50ns	-1μs	-100μs
Write Energy	-0.1nJ/b	<1nJ/b	0.1-1nJ/b
Read Latency	50ns	50 - 100 ns	10-25 μs
Read Energy	-0.1nJ/b	<<1nJ/b	<<1nJ/b
Byte-addressable	Yes	Yes	No

NVM- > Storage

1.5-8x





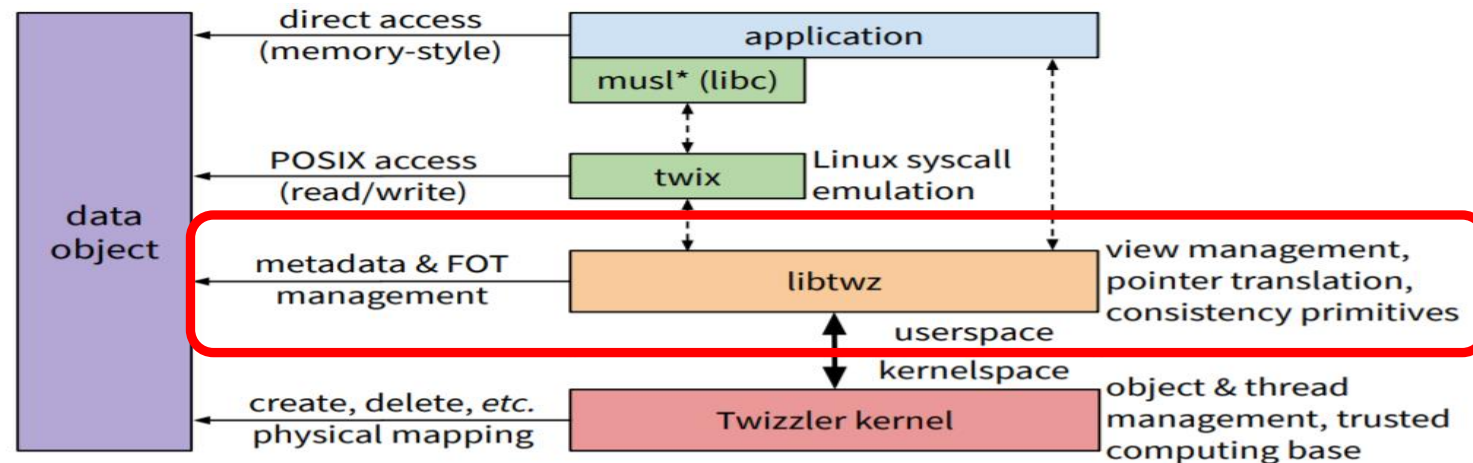
Why? Persistent data should be operated on *directly* and like *memory*

Requirements

- OS must evolve to support future trends in memory hierarchy organization.
- Twizzler is an OS designed around NVM for direct access to memory.
- Low overhead by removing the kernel from persistent data access path

Library OS

- ✓ Design
 - Object Management
 - Address space Management
 - Persistent Pointers
- ✓ Open-Source

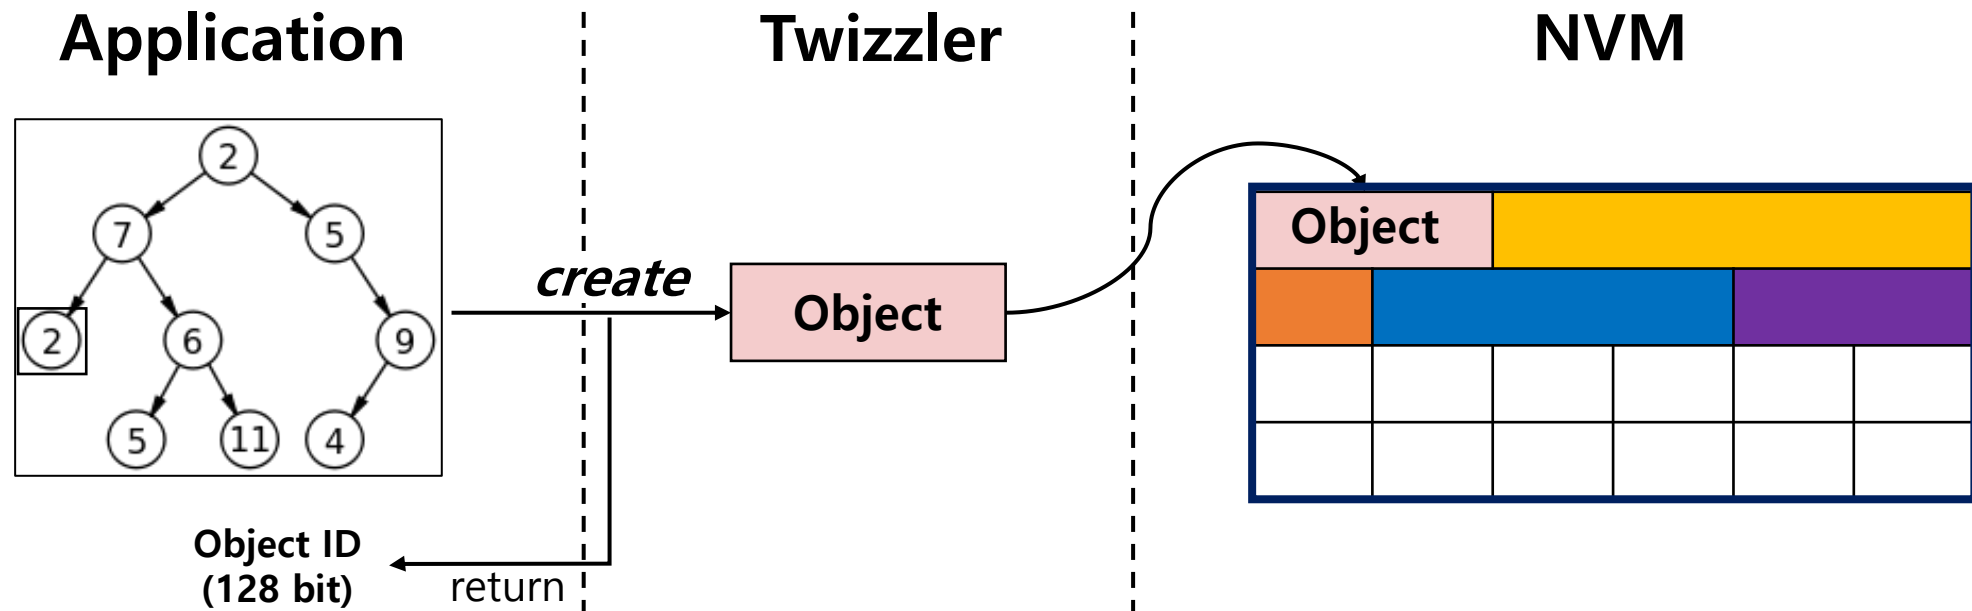


* modified must to change linux syscalls into function calls

Figure 1: Twizzler system overview. Applications link to `musl` (a C library), `twix` (our Linux syscall emulation library), and `libtwz` (our standard library).

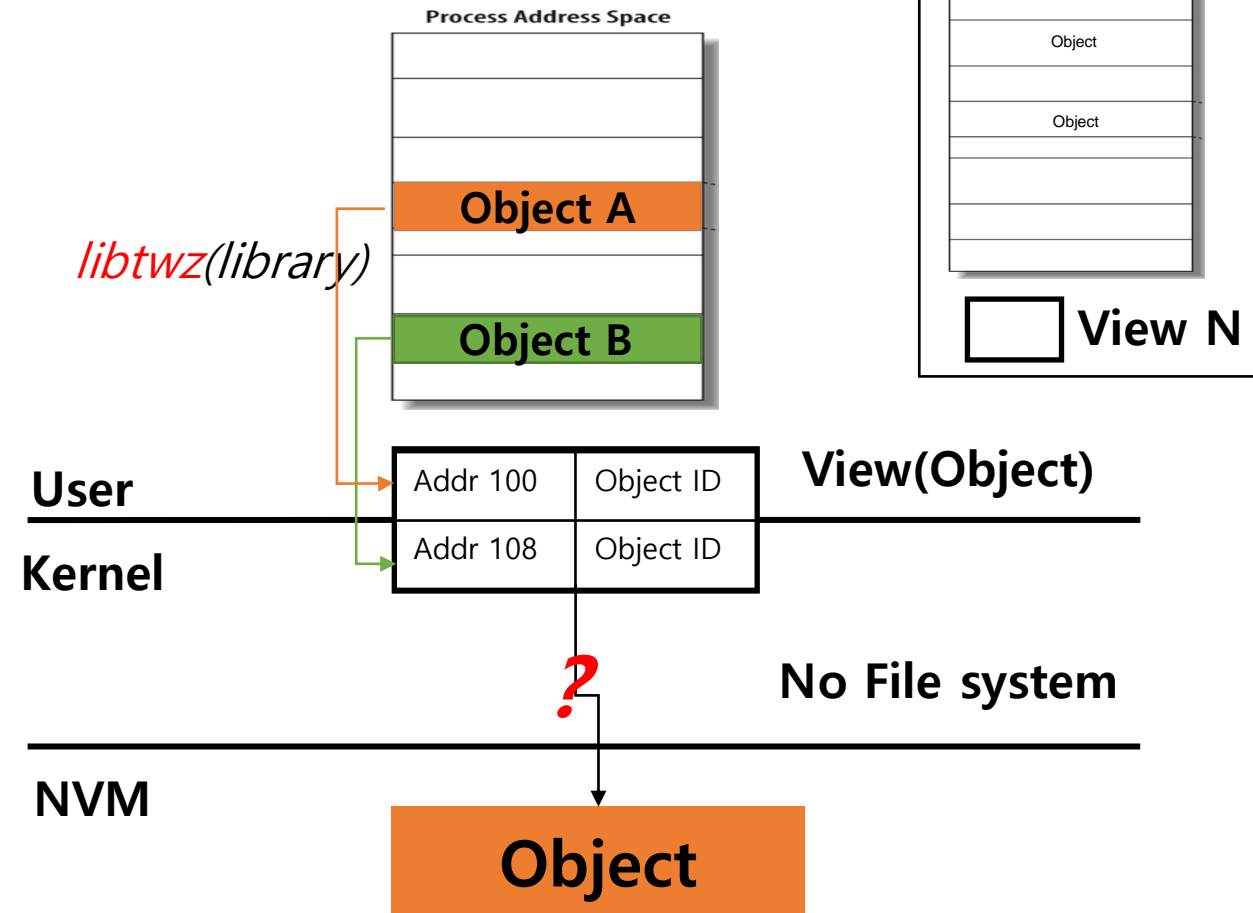
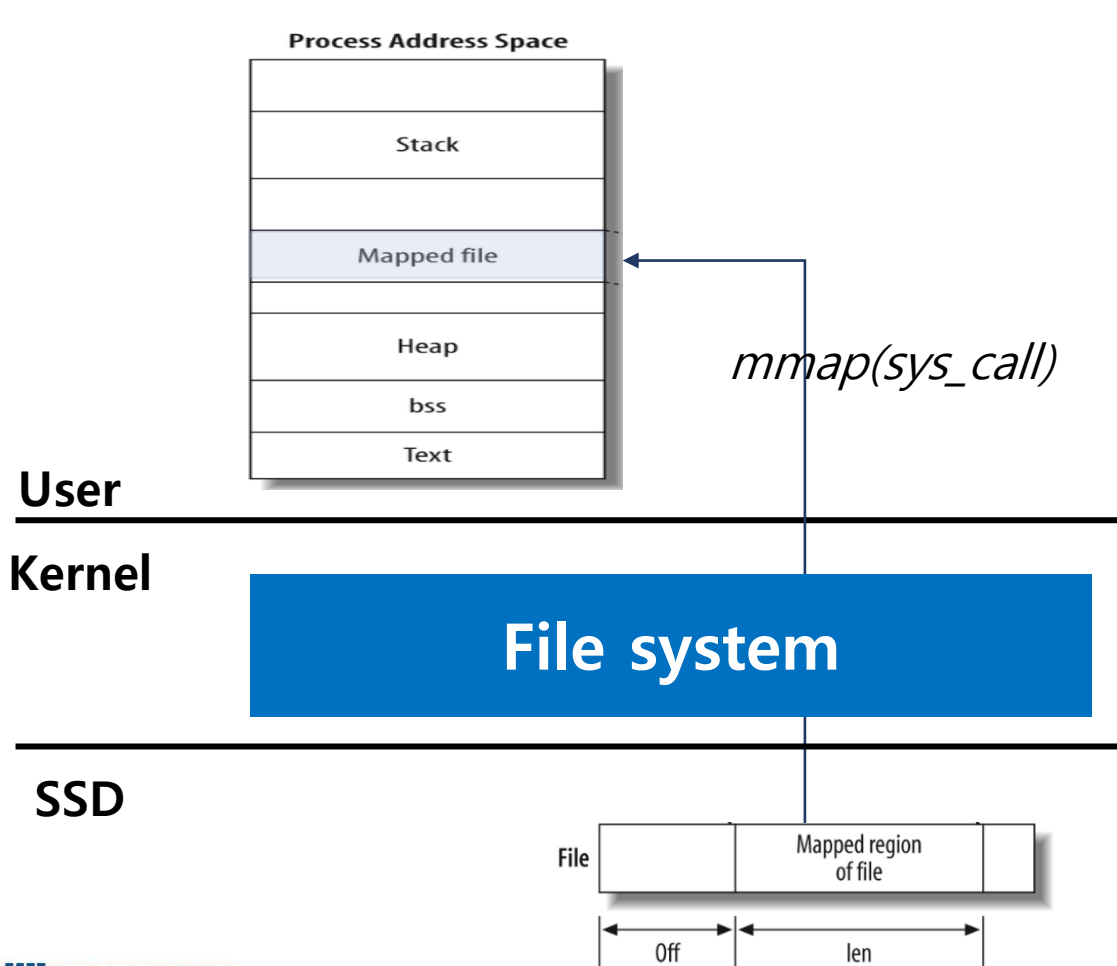
Object Management

- ✓ Twizzler organizes data into *objects*
- ✓ Twizzler uses objects as the unit of access control
 - Read/Write/Execute Permissions (MMU)
- ✓ An object is flexible in its contents (ex. A single node, the entire tree)
- ✓ Objects are created/deleted by the *create/delete* system call



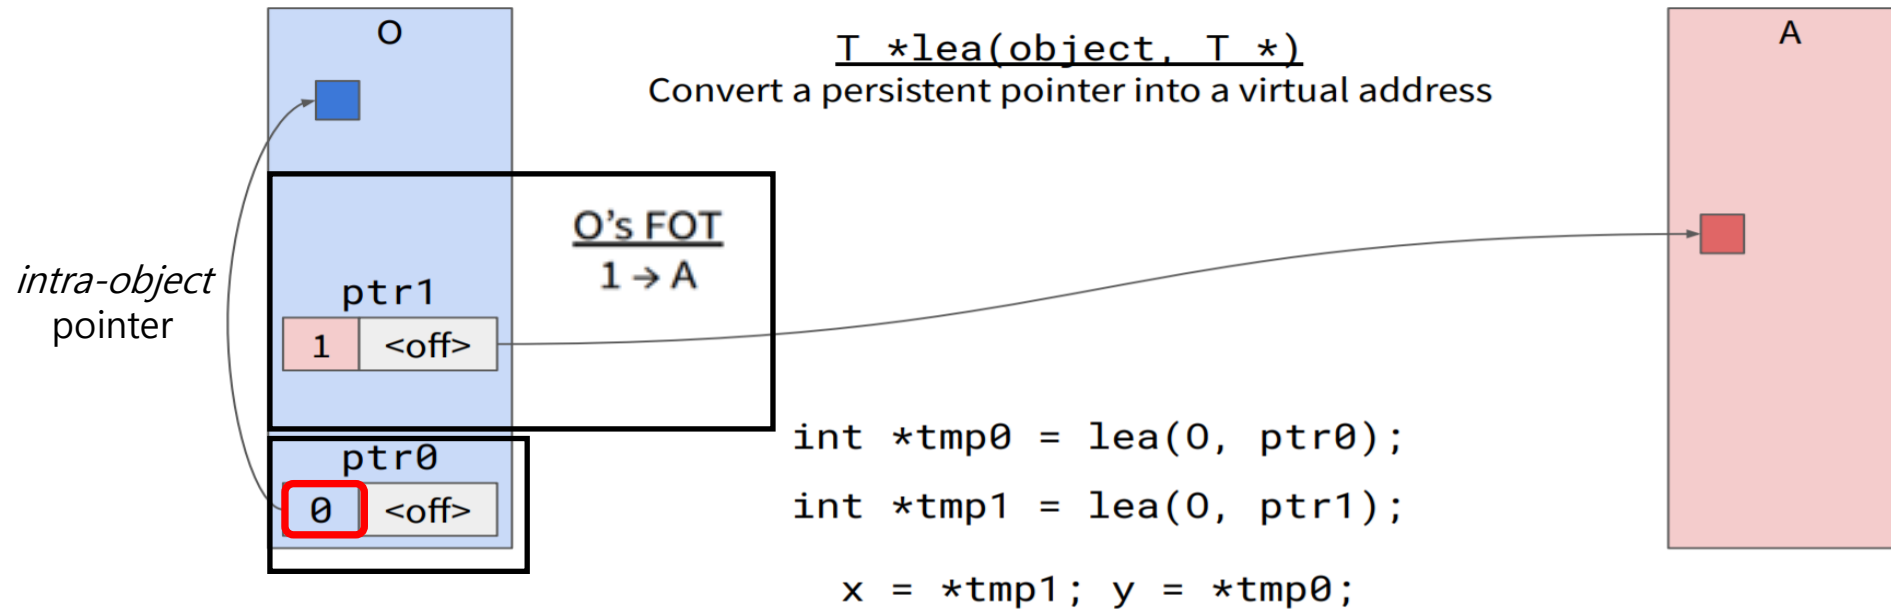
Address Space Management

- ✓ Twizzler maps persistent objects into the VA space via *libtwz*
- ✓ Kernel and user-space share a view that defines an AS layout



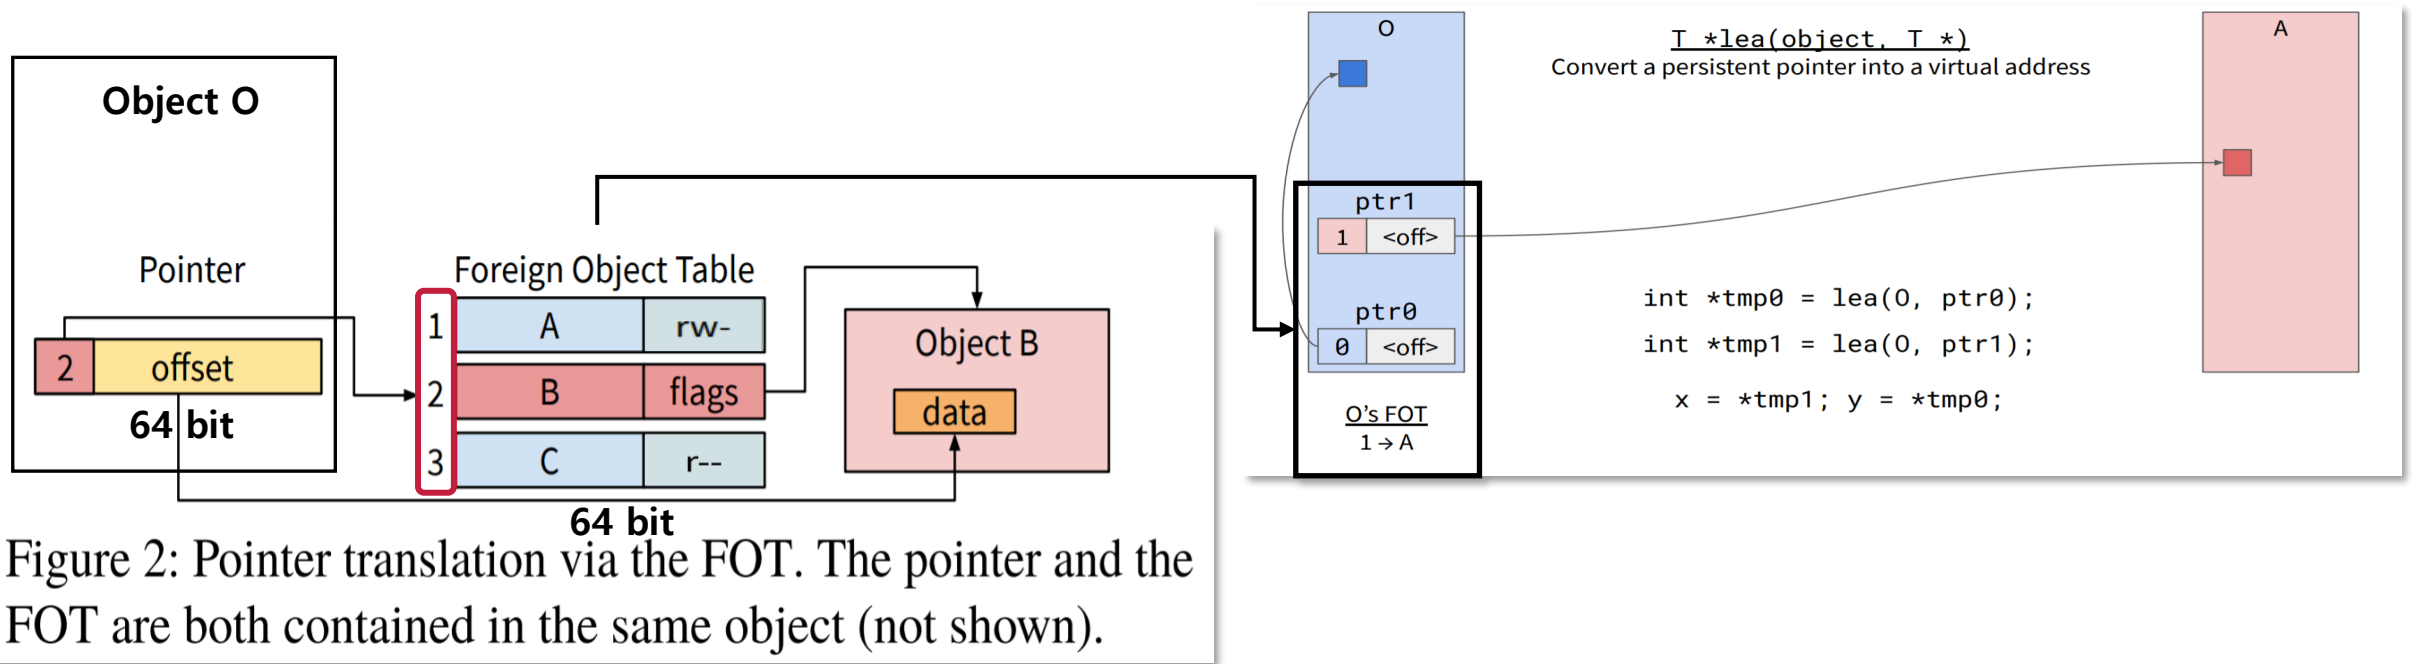
Persistent Pointers

- ✓ Twizzler provides **cross-object persistent pointers**
 - A cross-object pointer is stored as a 64 bit **FOT_idx:offset** value
- ✓ Twizzler use indirection through a per-object *foreign object table* (FOT)
 - The FOT is an array of entries that each stores an object ID and flags



Persistent Pointers

- ✓ Twizzler provides cross-object persistent pointers
 - A cross-object pointer is stored as a 64 bit FOT_idx:offset value
- ✓ Twizzler use **indirection** through a per-object *foreign object table (FOT)*
 - The FOT is an array of entries that each stores an object ID and flags



FOT entry of >0 means “cross-object”—points to a different object.

128bit Direct pointer → 64bit pointer → Cache Hit



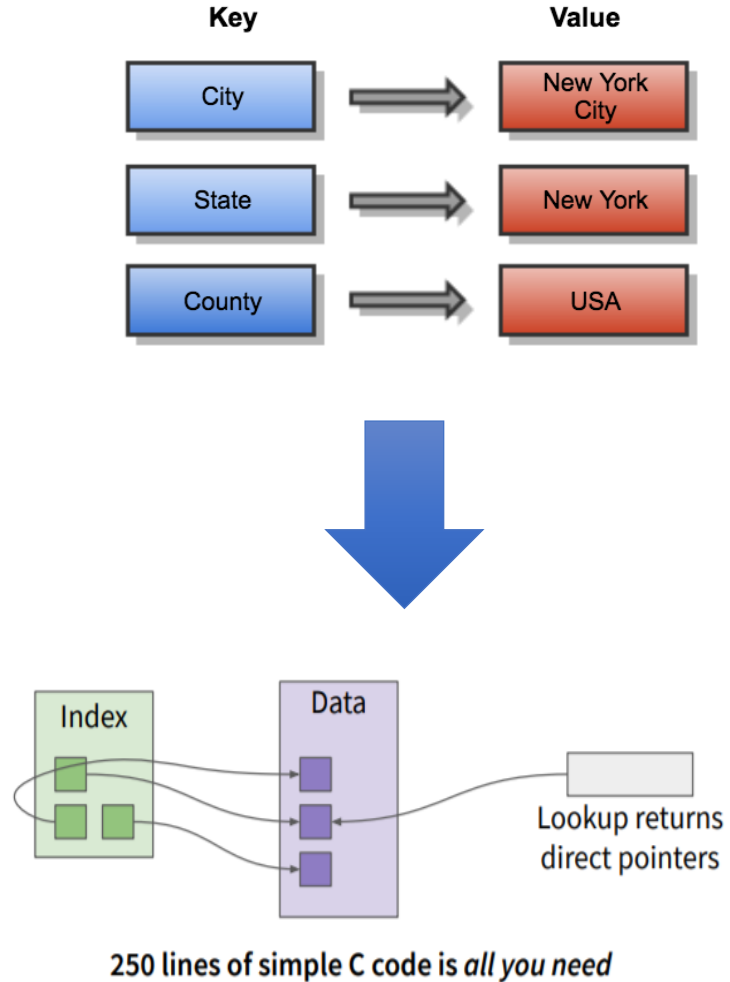
Micro benchmarks

Table 1: Latency of common Twizzler operations.

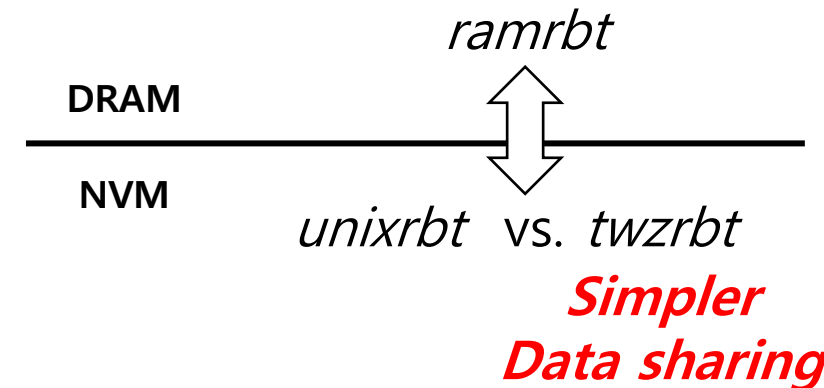
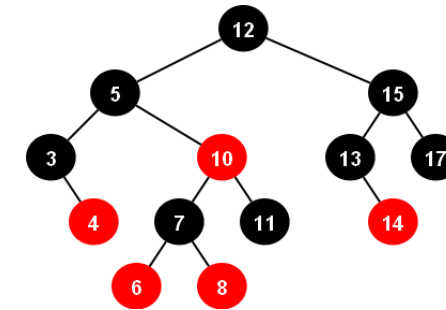
Pointer Resolution Action	Average Latency (ns)
Uncached FOT translation	27.9 ± 0.1
Cached FOT translation	3.2 ± 0.1
Intra-object translation	0.4 ± 0.1
Mapping object overhead	49.4 ± 0.2

Case Study

1) Key-Value Store



2) Red-Black Tree



Performance

1) Key-Value Store

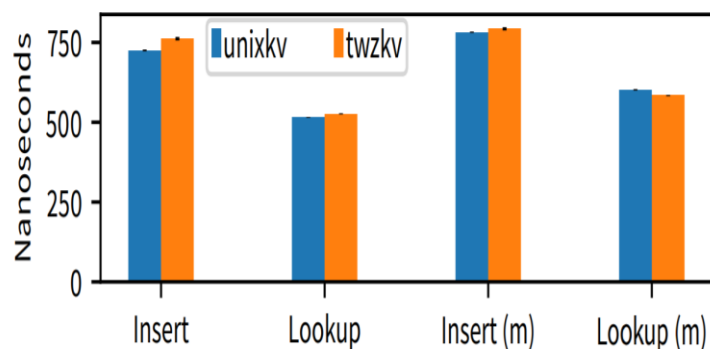
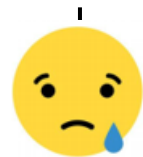


Figure 6: Latency of insert and lookup in twzkv and unixkv. An “(m)” indicates support for multiple data objects.

2) Red-Black Tree

We measured the latency of insert and lookup of 1 million 32-bit integers on both unixrbt and twzrbt. The insert and lookup latency of twzrbt was 528 ± 3 ns and 251.8 ± 0.5 ns, while insert and lookup latency of unixrbt was 515 ± 2 ns and 213 ± 1 ns. The modest overhead comes with significantly improved flexibility, as unixrbt does not support cross-object trees, and less support code (unixrbt manually implements mapping and pointer translations). Note that even though there is lookup overhead in twzrbt, this overhead did not predict the results of a larger program—the SQL-Twizzler port used this red-black tree, and saw performance benefits over block-based implementations.



3) SQLite

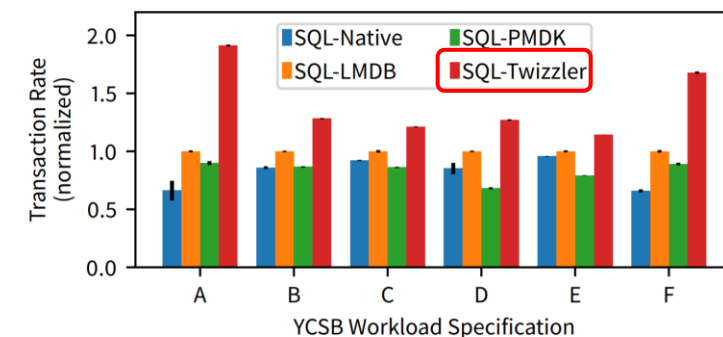


Figure 4: YCSB throughput normalized (higher is better).

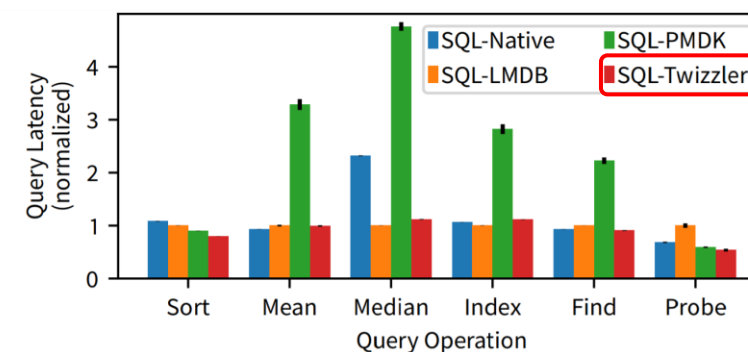


Figure 5: Query latency normalized (lower is better).

- **Future trends in memory hierarchy**
 - Byte-addressable, non-volatile memory NVM
- **Low latency and the direct-access nature of NVM**
 - Remove the kernel from I/O path, need not to serialize
 - Object/Address space management, persistent pointers
- **Evaluation**
 - Simple implementation, high performance



Q&A

2020. 08. 10

Presentation by Han, Yejin

32164881@dankook.ac.kr