

# LinnOS : Predictability on Unpredictable Flash Storage with a Light Neural Network

*Mingzhe Hao, Levent Toksoz, Nanqinqin Li, Edward Edberg Halim,  
In 2020 USENIX Symposium on Operating Systems Design and Implementation*

2021. 01. 26


Presentation by Han, Yejin

hyj0225@dankook.ac.kr

# Contents

1. Introduction
2. Background
3. Motivation
4. LinnOS
5. Evaluation
6. Conclusion






**Size of 2D & 3D NAND Flash Market, in USD billion, Global, 2015 – 2025**



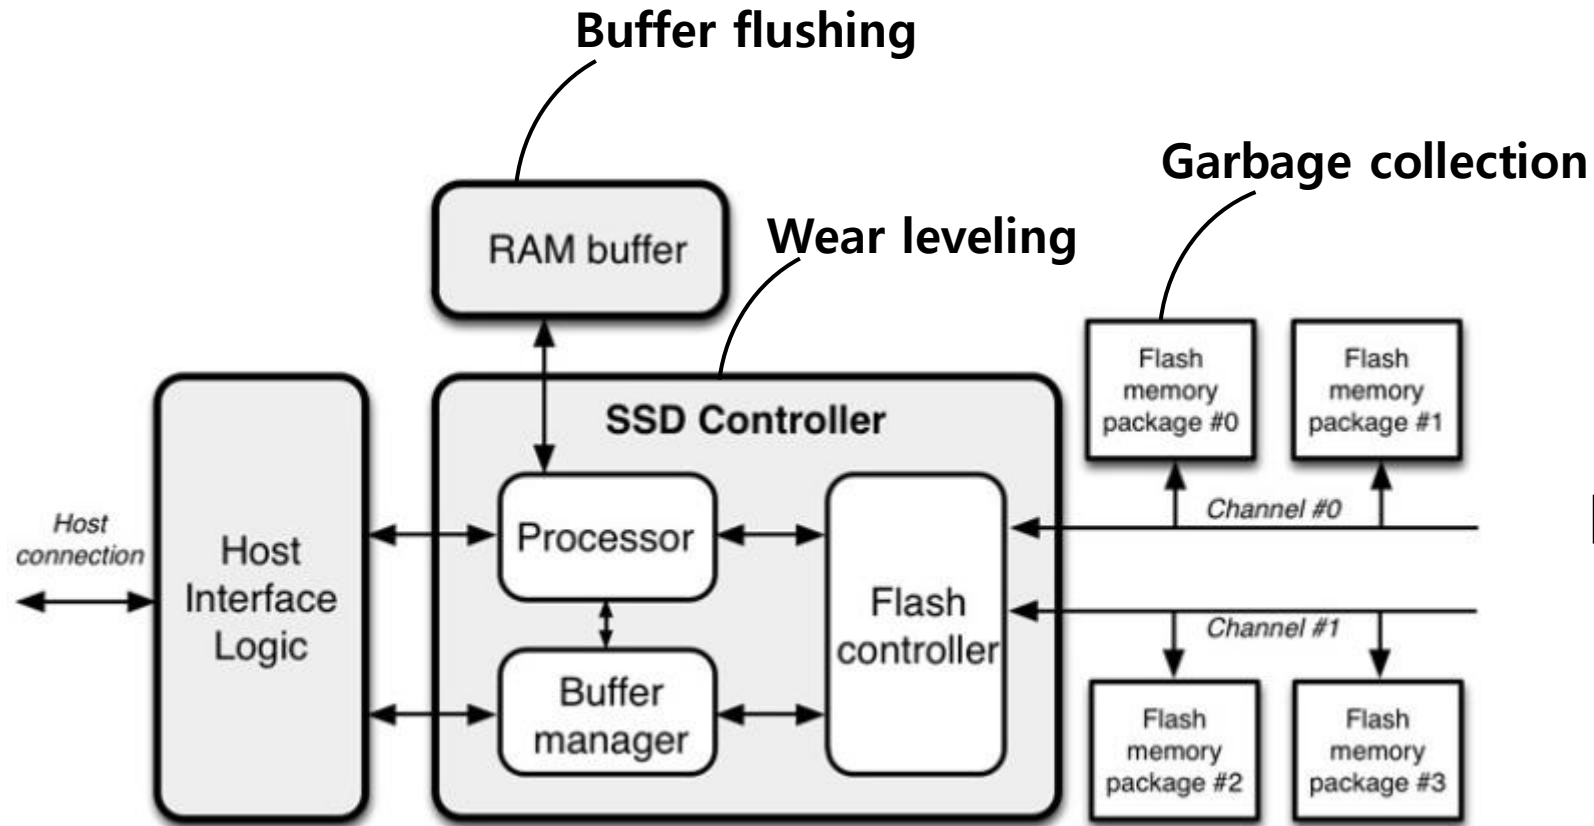
- Faster and faster SSDs are available
- SSD internal complexity grows




## Unpredictable latency

## Unpredictable latency

- SSD internal complexities threat to performance stability



### Unpredictable latency at individual-I/O level



# Conventional Approaches for predictability

## White/gray-box

- Re-architect device internals/interface

**KAML: A Flexible, High-Performance Key-Value SSD**

**AutoSSD: an Autonomic SSD Architecture**

**Tiny-Tail Flash: Near-Perfect Elimination of Garbage Collection Tail Latencies in NAND SSDs**

...



Powerful



Need to modify hardware

## Black-box

- SSD-aware filesystems and applications

**F2FS: A New File System for Flash Storage**

**Strata: A Cross Media File System**

**NOVA-Fortis: A Fault-Tolerant Non-Volatile Main Memory File System**

...



Not modifying hardware



Considerable re-design  
in software stack

# Conventional Approaches for predictability

Black-box: fs/storage applications, **speculative execution**



mongoDB



*cassandra*



Hadoop


...

# Conventional Approaches for predictability

Black-box: fs/storage applications, **speculative execution**



## Hedged requests

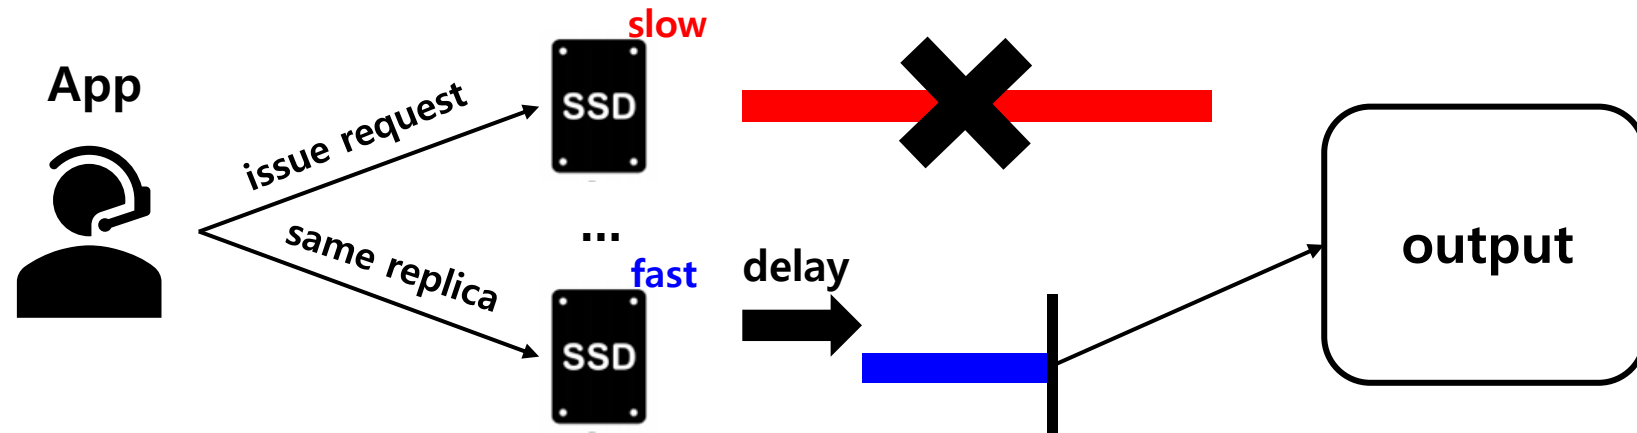


# Conventional Approaches for predictability

Black-box: fs/storage applications, **speculative execution**



## Hedged requests




**Mitigate every slow I/O!!**



## Device-Agnostic

### speculative execution


- Passively **wait** due to black-box



## Learn the device behavior

### LinnOS

- Proactively **infer** the black-box
- applications can know **in advance** whether performance expectations can be fulfilled



### LinnOS

- Machine learning for OS to learn black-box devices
- Admission control for clustered storage applications
- Per-I/O SSD performance prediction

## Usage scenario

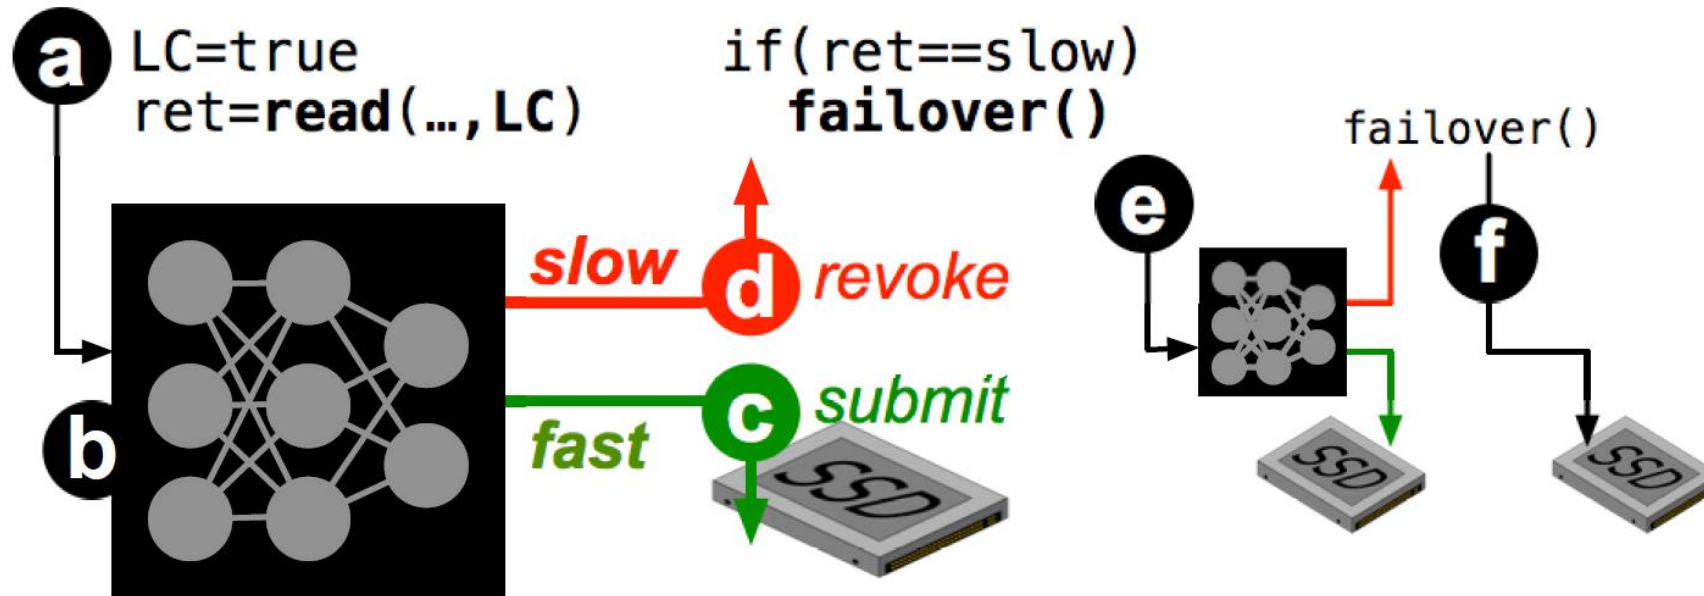


Figure 2: **Usage scenario.** *This usage scenario is explained in Section 3.1. “LC” implies latency critical.*

## Overall architecture


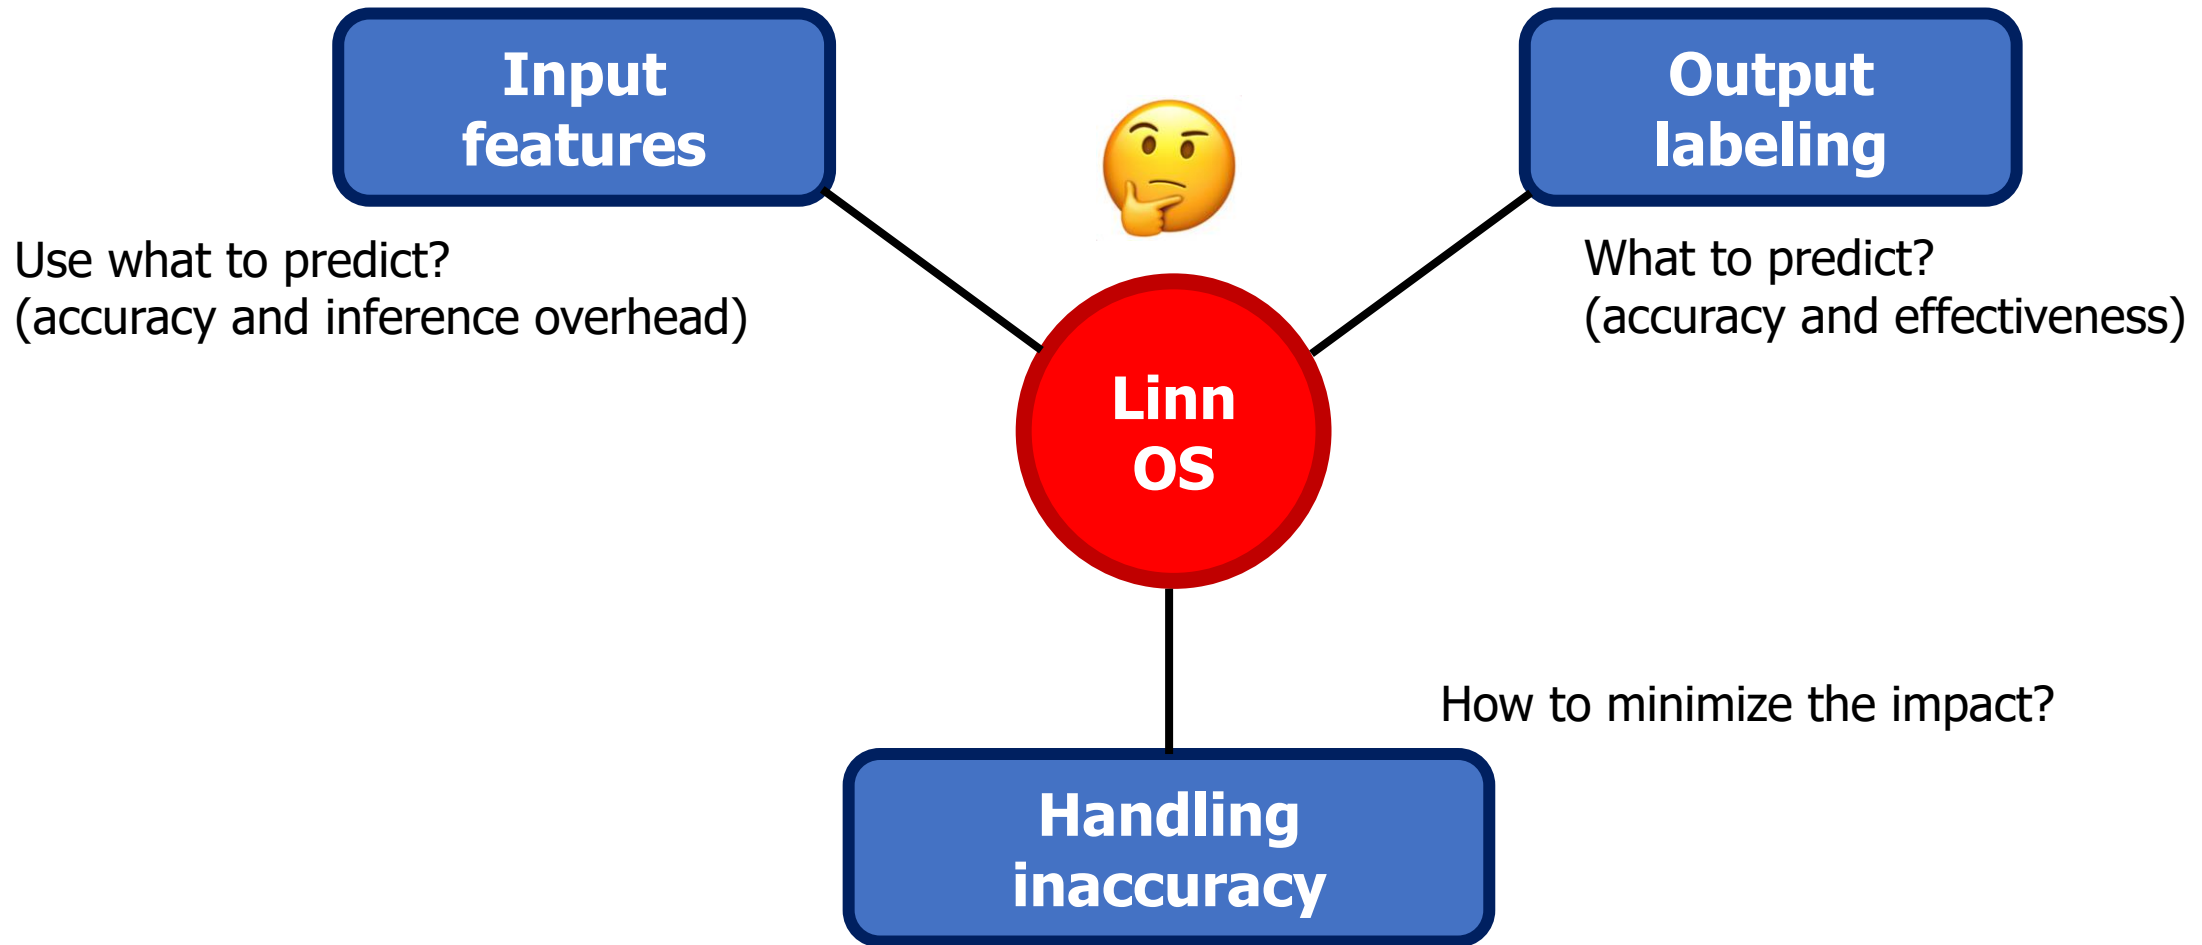


Figure 3: **LinnOS architecture.** The figure displays LinnOS architecture including LinnApp, as summarized in Section 3.2. The two SSD pictures represent the same SSD instance; the left one depicts tracing/training and the right one live inference on the SSD.

## Design Challenges



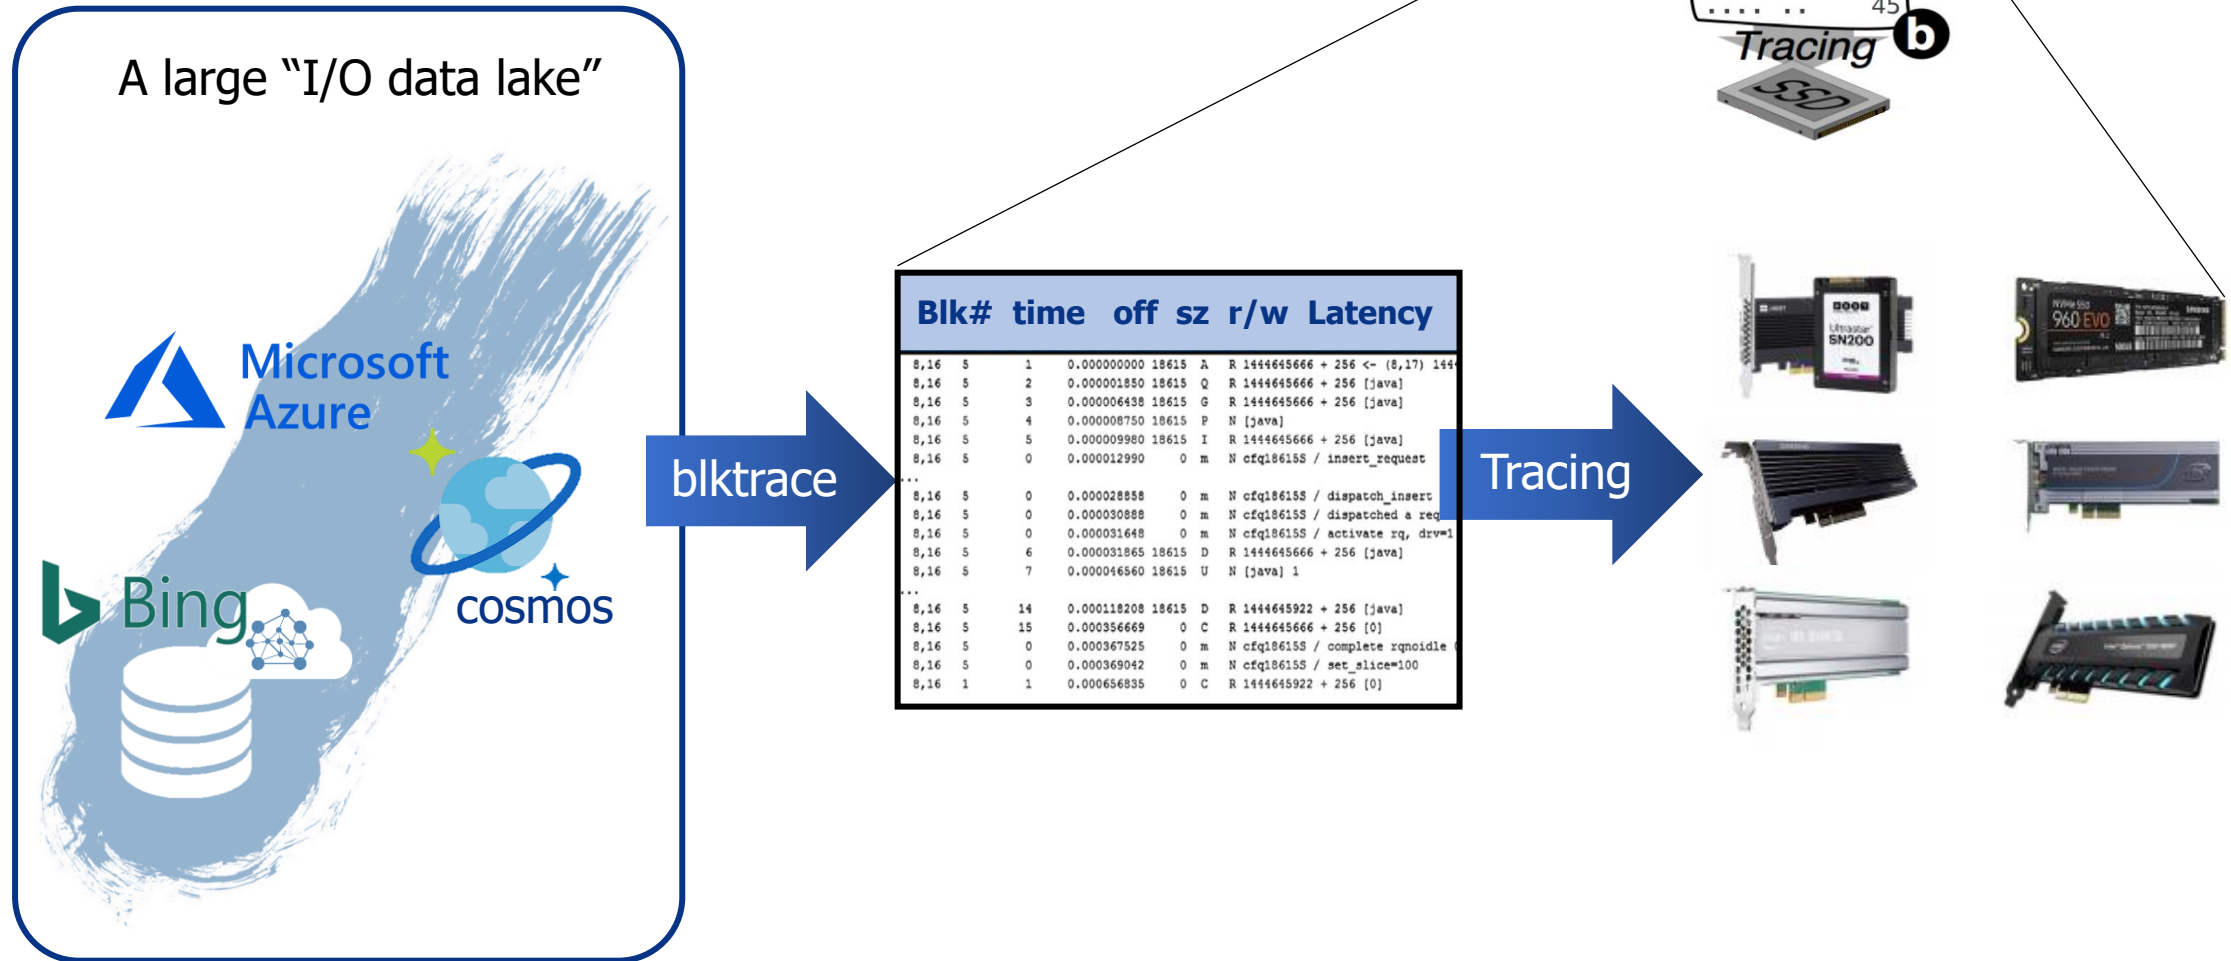


# Design Solutions

- 1) Training Data Collection
- 2) Labeling (with Inflection Point)
- 3) Light Nerual Network Model
- 4) Improving Accuracy
- 5) Improving Inference Time

## 1) Training Data Collection

- Online traces of real workload during busy-hour (e.g. midday)
- Submission time, block offset/size, read/write, latency per I/O



# Design Solutions

- 1) Training Data Collection
- 2) Output Labeling
- 3) Light Nerual Network Model
- 4) Improving Accuracy
- 5) Improving Inference Time

## 2) Output Labeling

**Ideal** Labeling

Exact  $\mu\text{s}$ -level latency  
(e.g.  $120\mu\text{s}$ ,  $70\mu\text{s}$ ...)



Flexible

Linear labeling  
(e.g.  $0-10\mu\text{s}$ ,  $10-20\mu\text{s}$ ...)



Too many labels

Hard to make accurate and fast

exponential labeling  
(e.g.  $2-4\mu\text{s}$ ,  $4-8\mu\text{s}$ ...)

**Only 60-70% accuracy** $64-128\mu\text{s}$  $128-256\mu\text{s}$ 

Mis-inferred

 $256-512\mu\text{s}$ 


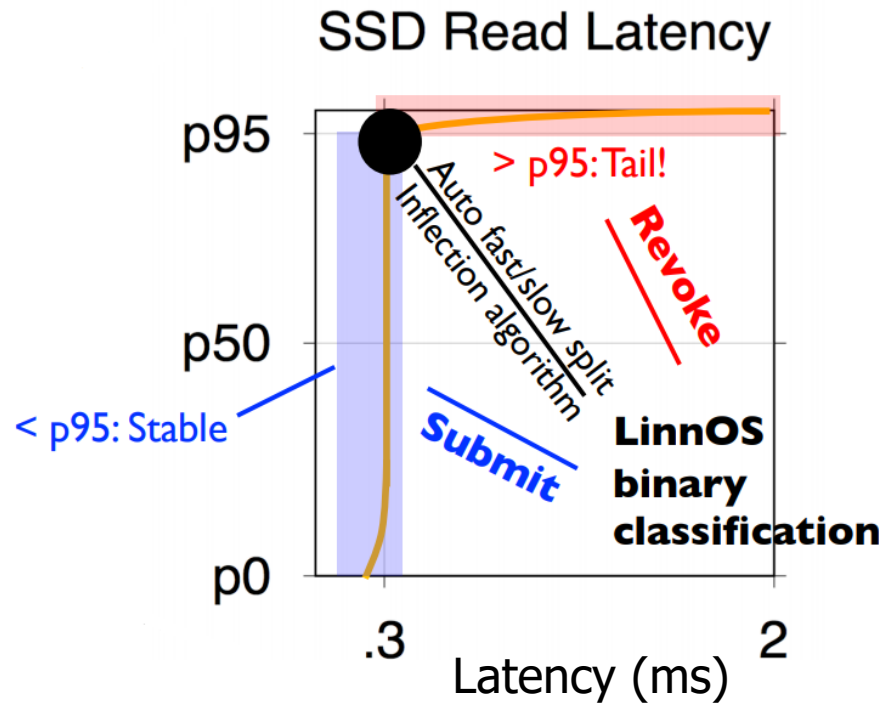
Truth

**Alternatives?**

## 2) Output Labeling

- Latencies often form a Pareto distribution with a high alpha number
- Users only worry about the tail behavior, not the precise latency

### Labeling with **inflection point**

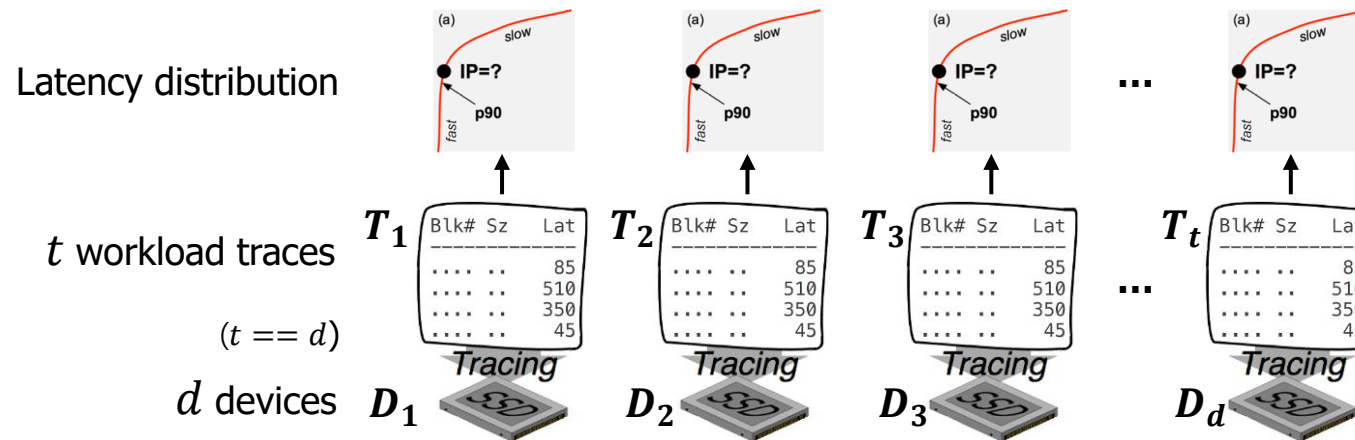


**How to find?**



## 2) Output Labeling

- Inflection Point Algorithm

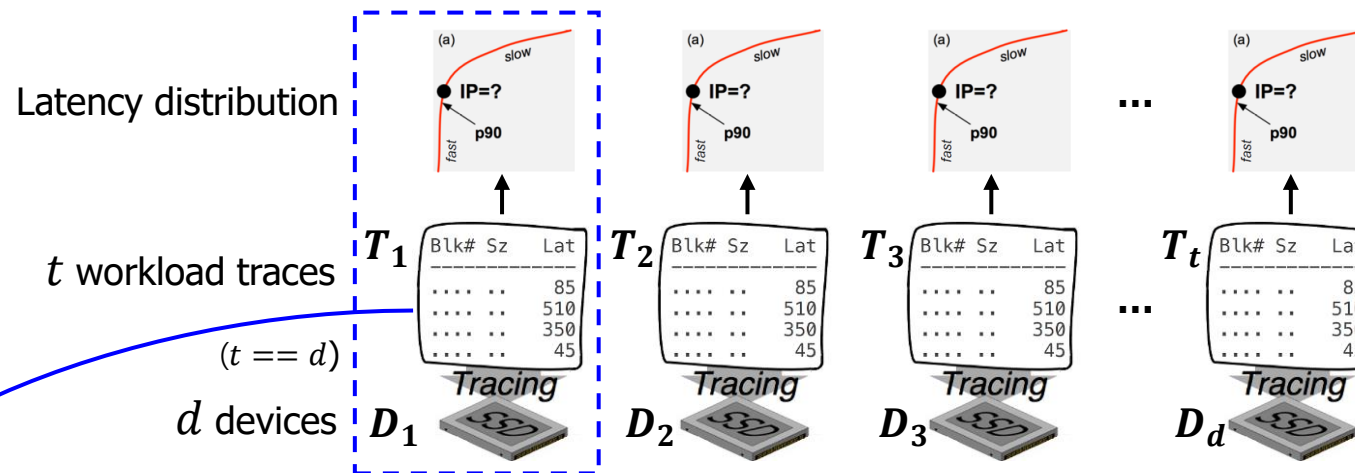


- Inflection Point Algorithm

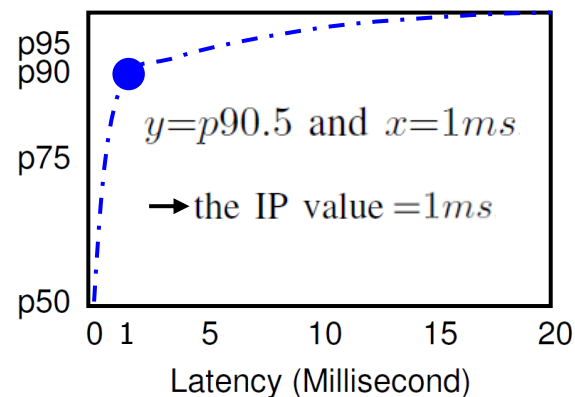


## 2) Output Labeling

- Inflection Point Algorithm



### 1) Pick a starting IP value




### 2) Simulate/repeat admission control

I/O request  $r_i$ 's latency value  $\leq 1ms$

→ the  $r_i$ 's new latency is set to be the same;

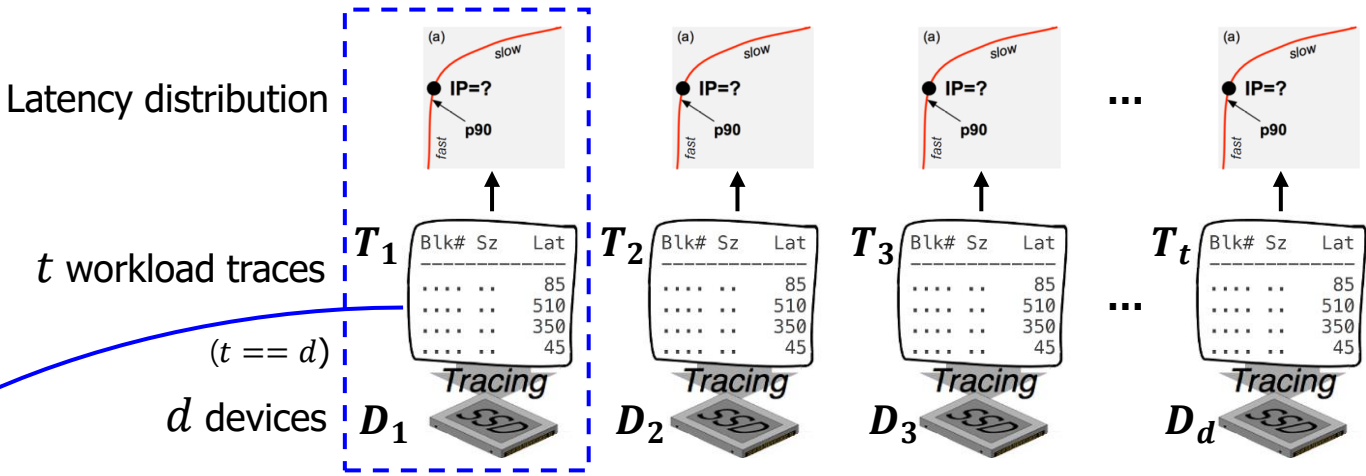
I/O request  $r_i$ 's latency value  $> 1ms$

→ revoked and failed over to randomly selected node  
selected node (e.g.,  $D_4$ )




2) Output Labeling

- Inflection Point Algorithm



1) Pick a starting IP value




2) Simulate/repeat admission control

I/O request  $r_i$ 's latency value  $\leq 1ms$

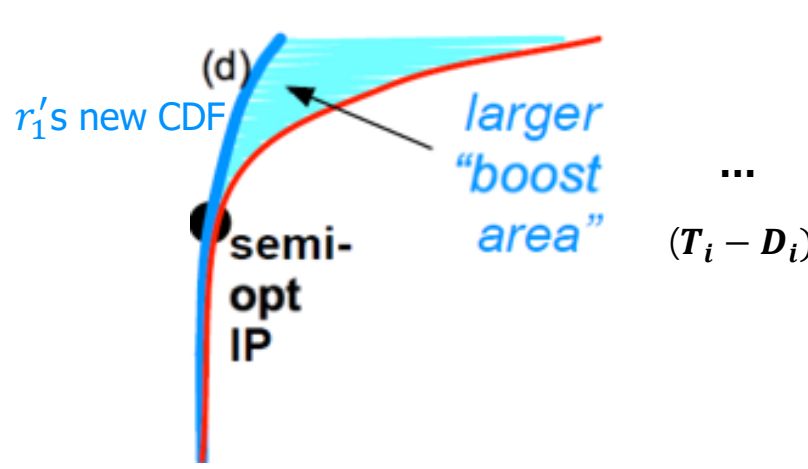
→ the new CDFs and pick the  $IP^{max}$

I/O request  $r_i$ 's latency value  $> 1ms$

→ revoked and failed over to randomly selected node selected node (e.g.,  $D_4$ )



3) New CDFs and pick the  $IP^{max}$



# Design Solutions

- 1) Training Data Collection
- 2) Labeling (with Inflection Point)
- 3) Input features
- 4) Improving Accuracy
- 5) Improving Inference Time



### 3) Input features

- The number of pending I/Os (4KB pages) when an incoming I/O arrives
- The latency of the 4 most-recently completed I/Os
- The number of pending I/Os when each of the 4 completed I/Os arrived




Figure 6: **Light neural network.** The figure depicts LinnOS 3-layer neural network explained in Section 4.3.

### 3) Input features

- The number of pending I/Os (4KB pages) when an incoming I/O arrives
- The latency of the 4 most-recently completed I/Os
- The number of pending I/Os when each of the 4 completed I/Os arrived




Figure 6: **Light neural network.** The figure depicts LinnOS 3-layer neural network explained in Section 4.3.

### 3) Input features

- The number of pending I/Os (4KB pages) when an incoming I/O arrives
- The latency of the 4 most-recently completed I/Os
- The number of pending I/Os when each of the 4 completed I/Os arrived




Figure 6: **Light neural network.** The figure depicts LinnOS 3-layer neural network explained in Section 4.3.

### 3) Input features

- The number of pending I/Os (4KB pages) when an incoming I/O arrives
- The latency of the 4 most-recently completed I/Os
- The number of pending I/Os when each of the 4 completed I/Os arrived

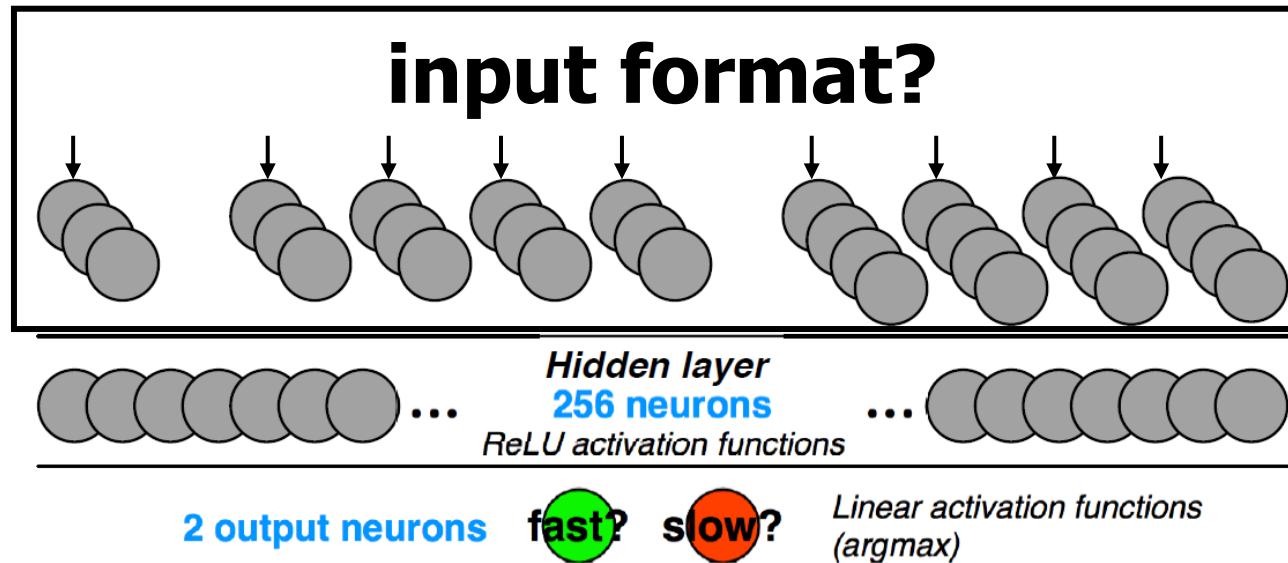
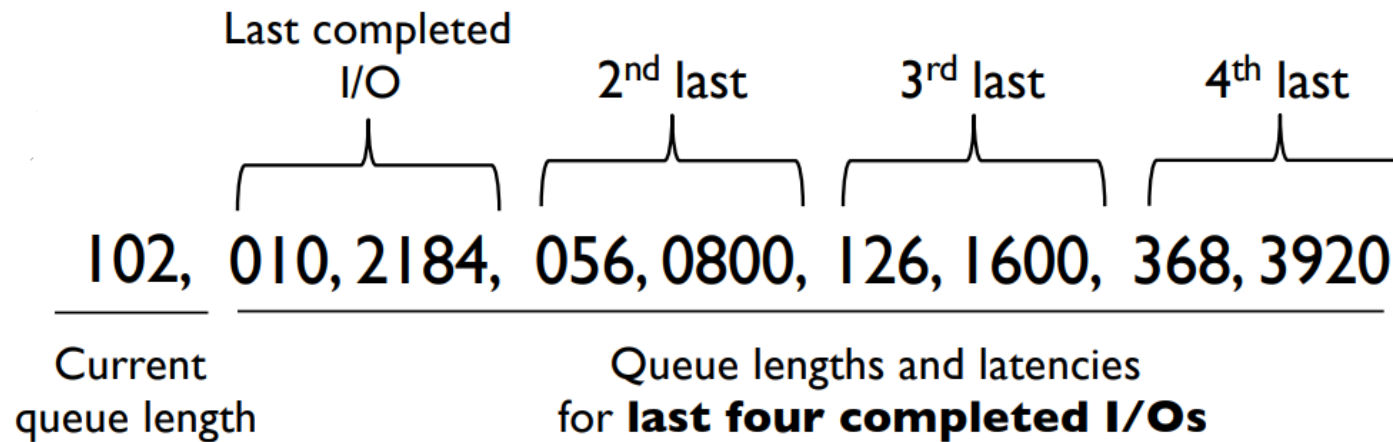


Figure 6: **Light neural network.** The figure depicts LinnOS 3-layer neural network explained in Section 4.3.

### 3) Input features

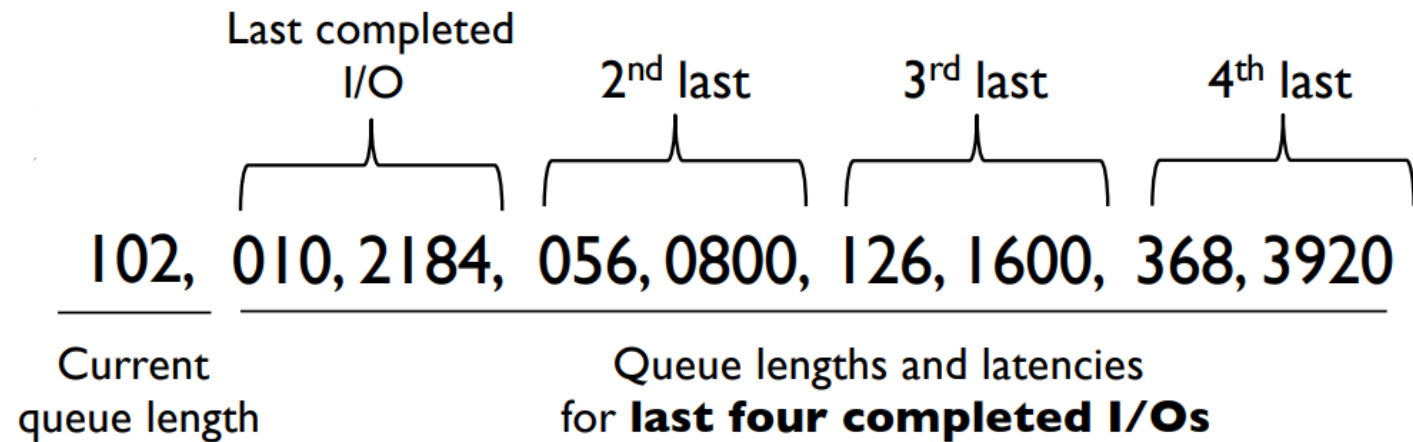
- Format the number of pending I/Os into three decimal digits
- Format  $\mu$ s latency value into four digits
- LinnOS model takes 31 input features, each a one-digit decimal number





### 3) Input features

- Format the number of pending I/Os into three decimal digits
- Format  $\mu$ s latency value into four digits
- LinnOS model takes 31 input features, each a one-digit decimal number



Split into individual digits


1,0,2, 0,1,0, 2,1,8,4, 0,5,6, 0,8,0,0, 1,2,6,1,6,0,0, 3,6,8,3,9,2,0

31  
features

### 3) Input features

- Format the number of pending I/Os into three decimal digits
- Format  $\mu$ s latency value into four digits
- LinnOS model takes 31 input features, each a one-digit decimal number

#### 3 fully-connected layers (31 - 256 - 2)

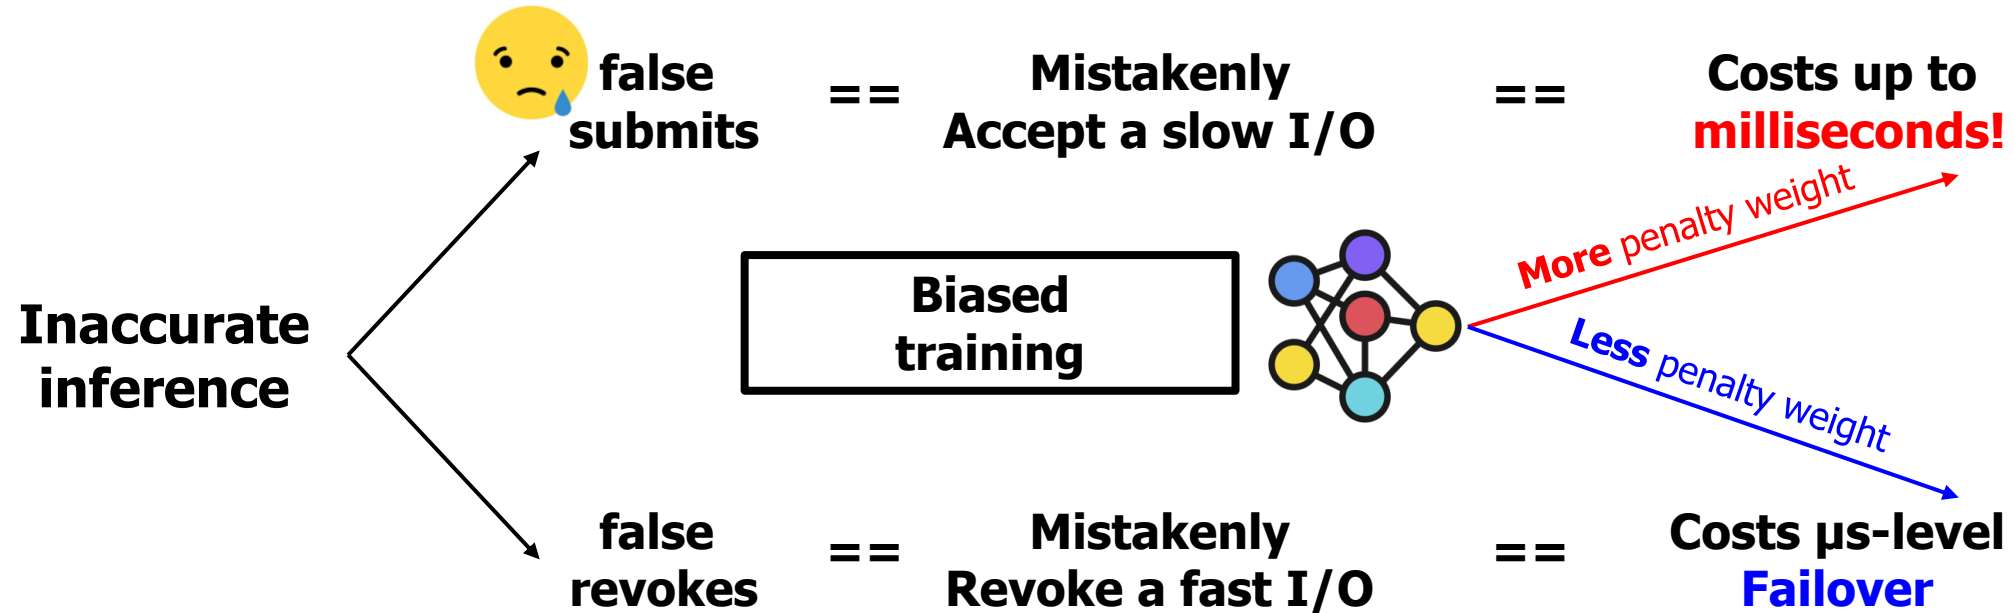


# Design Solutions

- 1) Training Data Collection
- 2) Labeling (with Inflection Point)
- 3) Light Nerual Network Model
- 4) Improving Accuracy
- 5) Improving Inference Time

#### 4) Improving Accuracy

- Wrong inference penalty is small for false revokes but high for false submits
- Use customized loss function: *biased training*
- Reduce false submits by allowing more false revokes



# Design Solutions

- 1) Training Data Collection
- 2) Labeling (with Inflection Point)
- 3) Light Nerual Network Model
- 4) Improving Accuracy
- 5) Improving Inference Time

### 5) Improving Inference Time

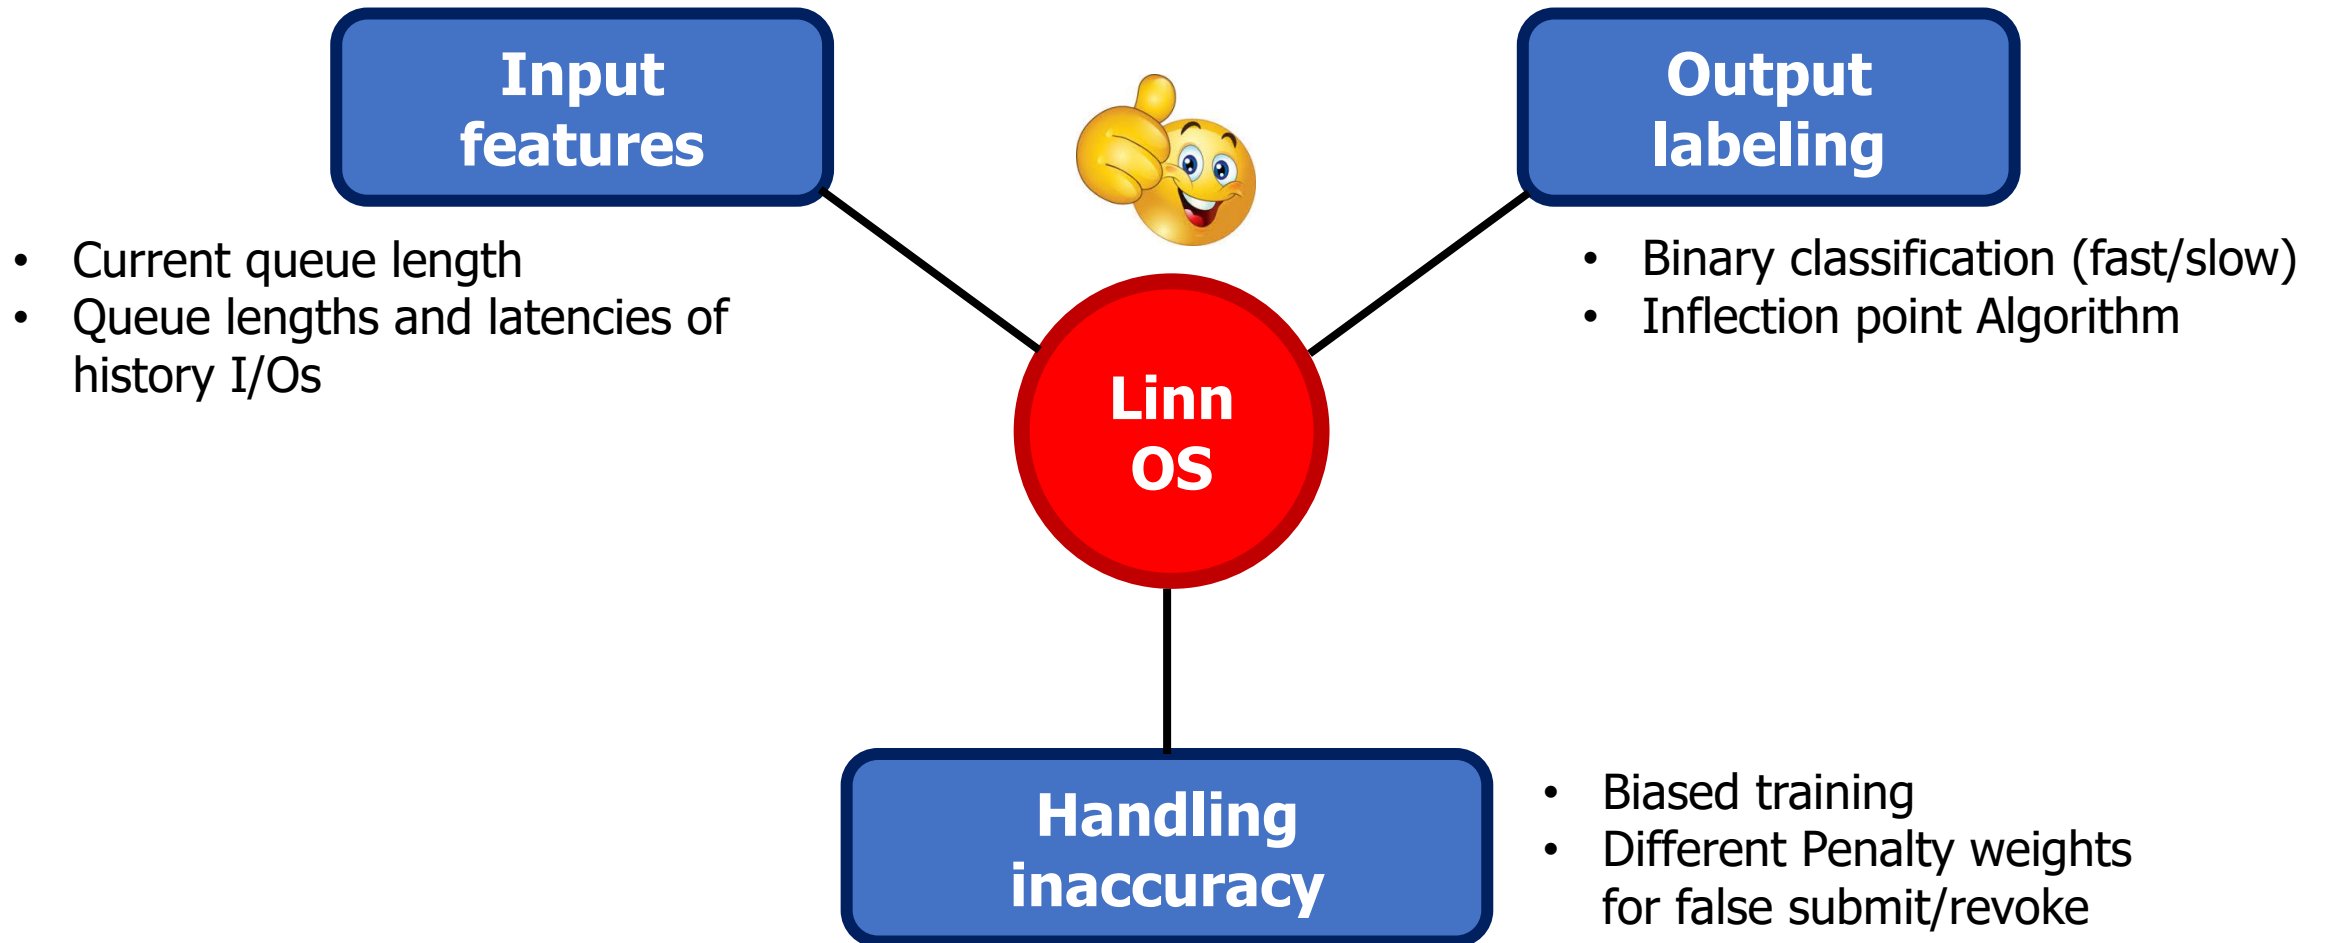
#### Quantization

- Storage functionalities are striping and partitioning using mod operations over integers
- Floating point calculations are expensive and hard to manage inside the OS
- Adopt DNN quantization by maintaining precision of three decimal points

#### Co-processors

- Can Utilize one additional CPU core (if available)
- reduce the inference time from 6 to 4 $\mu$ s with 2-threaded optimized matrix multiplication

## Summary




### Setup

#### 1) workload


real production SSD-level traces

- Microsoft Azure server
- Bing Index server
- Bing Select server
- Cosmos server




#### 2) devices

- 3 Samsung SM951 M.2 NVMe



Three **homogeneous consumer-level** SSDs

- Intel P4600
- Samsung PM1725a NVMe
- WD ultrastar DC SN200 NVMe




**OR** Three **heterogenous enterprise-level** SSDs

- 18 CPU core (36-thread)
- 128GB DRAM
- + 6 SSDs for accuracy evaluation

#### 3) The experiments

Methods compared:

- Baseline
- Cloning
- Hedging at p95
- Hedging at IP
- Simple heuristic
- Advanced heuristic
- LinnOS without hedging
- LinnOS





### ▪ Inflection Point (IP) Stability

	Consumer	Enterprise
Azure	p73.3, p77.0, p91.4	p91.0, p93.2, p97.8
BingIndex	p80.0, p94.5, p98.5	p80.1, p83.3, p97.0
BingSelect	p72.0, p76.9, p87.2	p75.3, p83.7, p86.8
Cosmos	p73.4, p82.5, p84.1	p83.2, p84.8, p95.1

Table 2: **Inflection point (IP) settings.** This table, as explained in Section 5.2, shows the IP values that our algorithm in Section 4.2.1 computed for every workload-device pair.




Figure 7: **IP stability.**

- reduces average latency by 9.6-79.6%

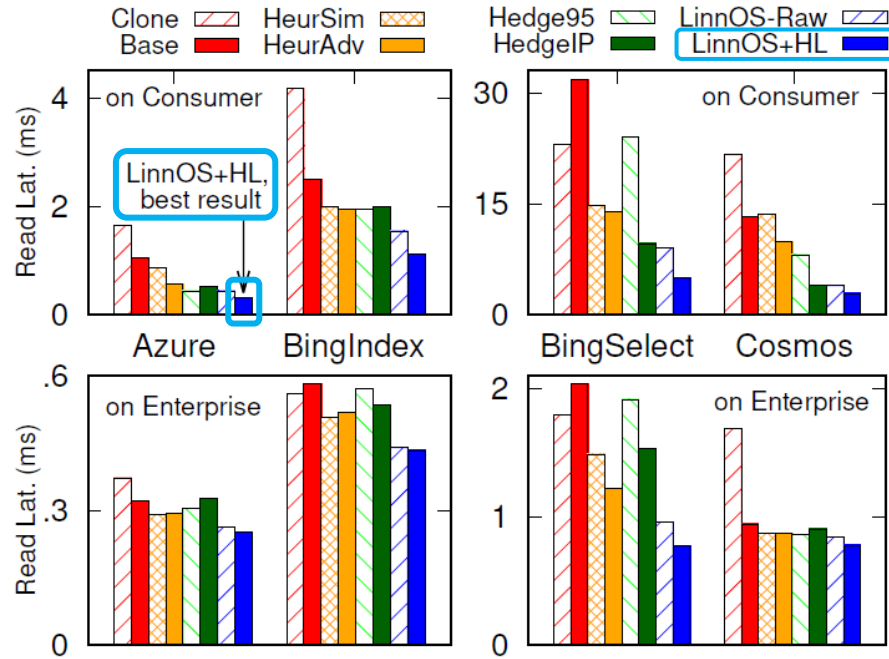
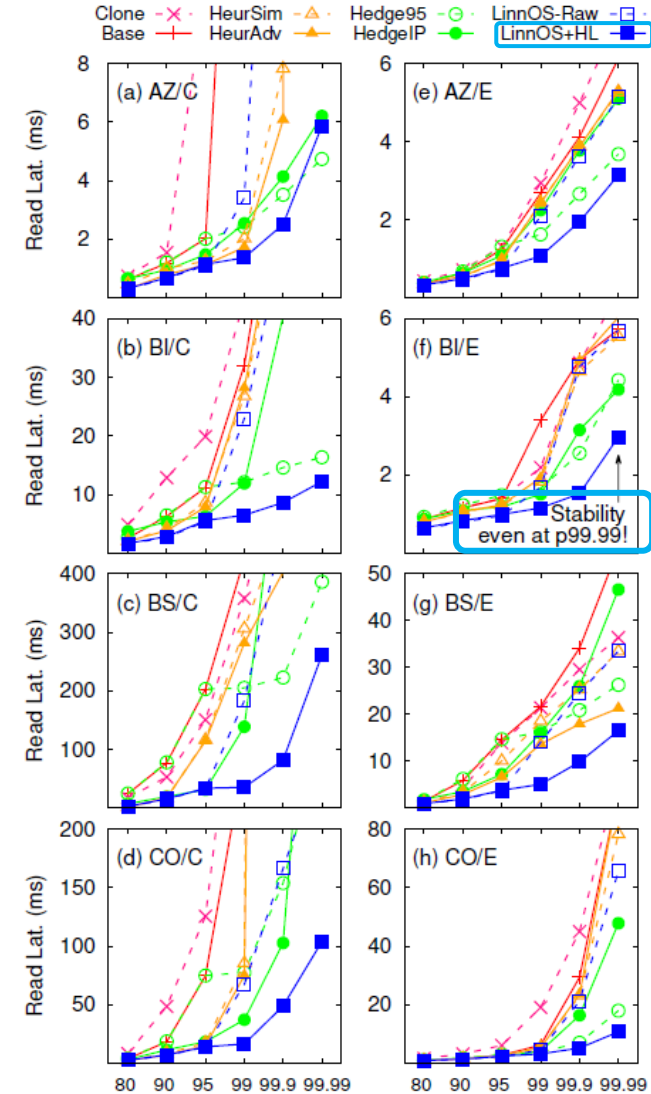


Figure 8: **Average latencies.** The figures show that LinnOS consistently outperforms all other methods, as explained in Section 5.3. The top and bottom graphs represent experiments on the consumer and enterprise arrays, respectively.

- Brings stable latencies



- high accuracy with low false submit/revoke

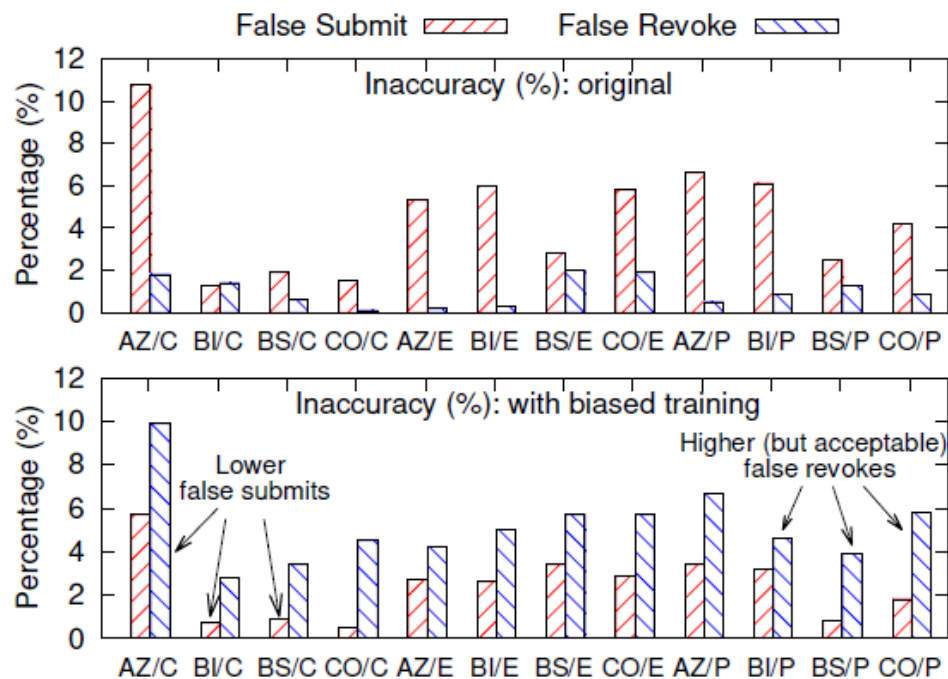
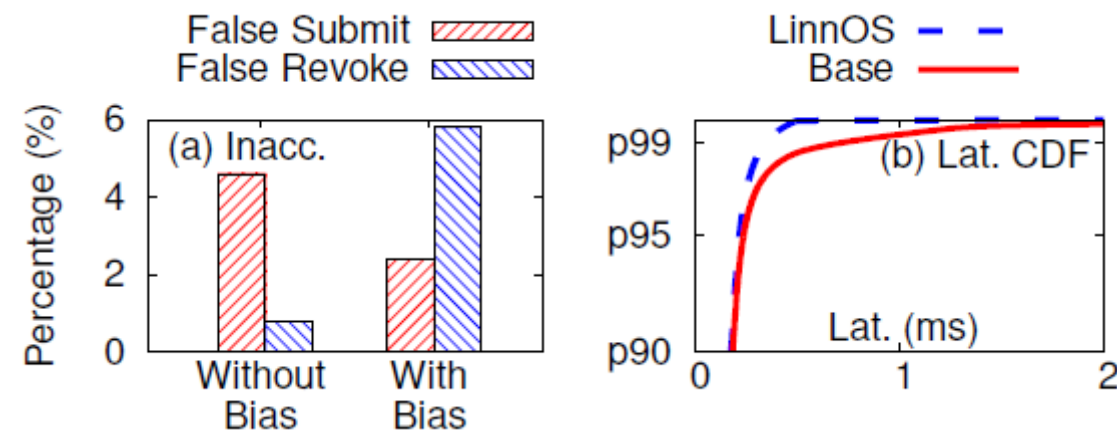


Figure 10: **Low inaccuracy.** The figure shows the percentage of false submits and false revokes. Note that only false submits really matter (see Section 5.4). Additionally, “P” represents other device models that we can access from a public cloud. For graph readability, here for “P” we only show the results for one device model, while the observations stand across the rest. In total, the accuracy evaluation covers 10 device models (1C+3E+6P).

- Works on public traces



- Helps MongoDB and FS




Figure 13: **MongoDB on different filesystems.** This figure shows that LinnOS can easily help data applications achieve more predictable latency (§5.6.3).

### LinnOS

- Demonstrate that it is possible to incorporate machine learning into OS
- Outperforms many other popular methods
- Successfully predict the speed of every I/O to flash storage

# LinnOS : Predictability on Unpredictable Flash Storage with a Light Neural Network

*Mingzhe Hao, Levent Toksoz, Nanqinqin Li, Edward Edberg Halim,  
In 2020 USENIX Symposium on Operating Systems Design and Implementation*

## Thank You!

2021. 01. 26

Presentation by Han, Yejin

hyj0225@dankook.ac.kr