

Application-Managed Flash

Sungjin Lee, Ming Liu, Sangwoo Jun, and Shuotao Xu, Jihong Kim, Arvind
14th USENIX Conference on File and Storage Technologies (FAST '16). 2016.

2021. 06. 02

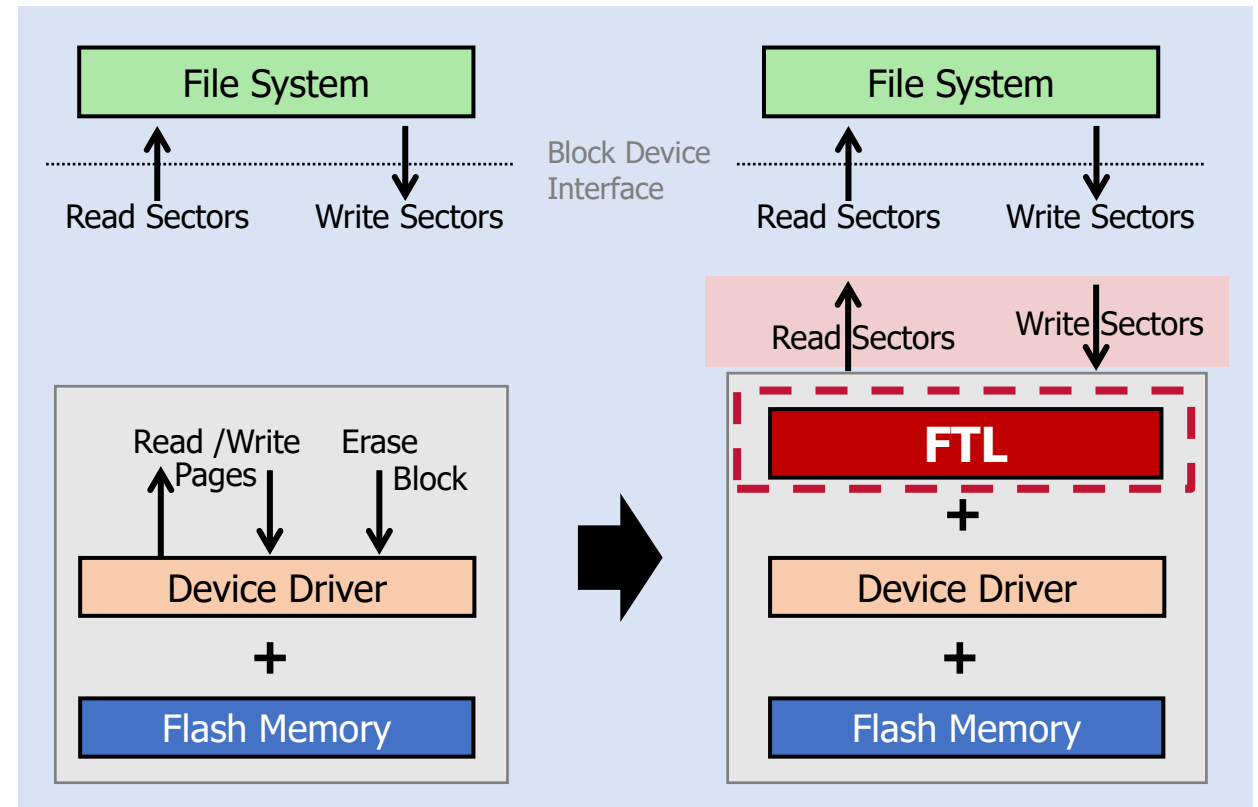

Presentation by Han, Yejin

hyj0225@dankook.ac.kr

Contents

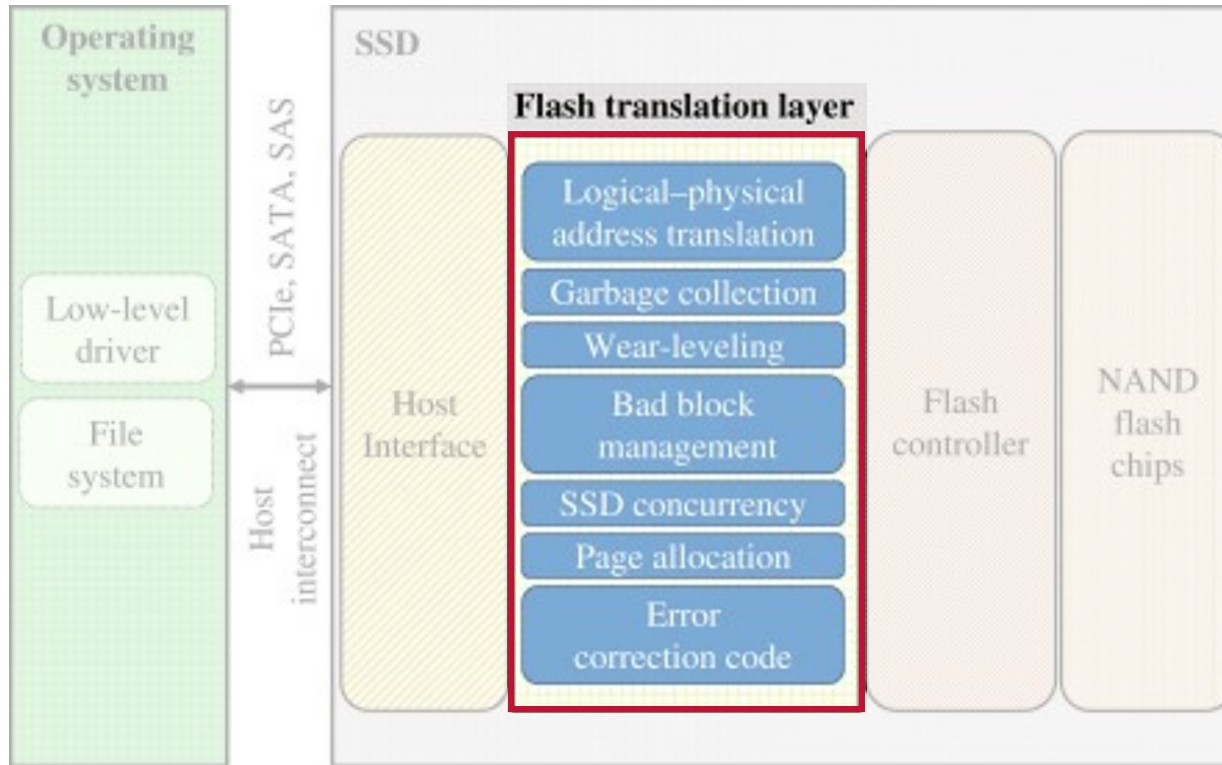
1. Introduction
2. Motivation
3. AMF
 - 3-1. Block I/O interface
 - 3-2. ALFS
 - 3-3. AFTL
4. Evaluation
5. Conclusion

- NAND flash SSD become the preferred storage device in consumer electronics and data centers
- FTL provide an block I/O abstraction for interoperability with HDDs



What are the purposes of the FTL?

- Managing address translation, Garbage collection, wear-leveling, ...



- 1) Require significant hardware resources
- 2) Extra I/Os for flash management (GC)
- 3) Host applications cannot predict the behavior of flash storage

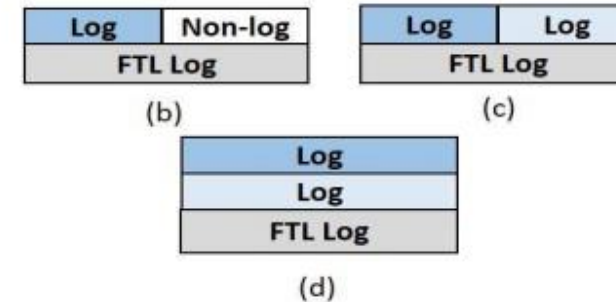
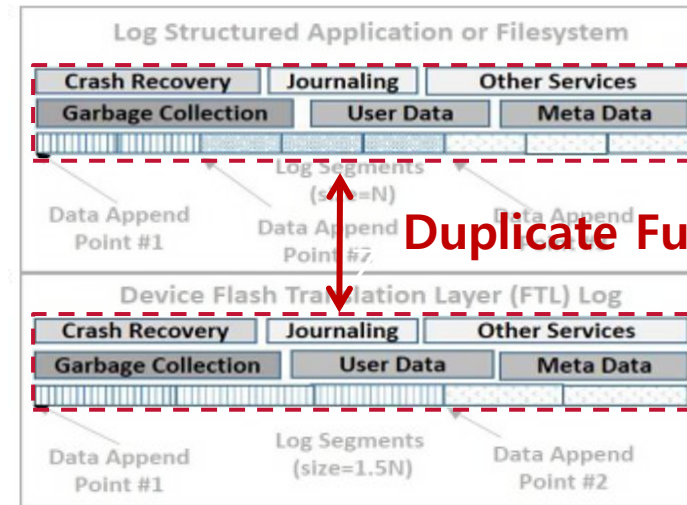


What are more issues of the FTL?

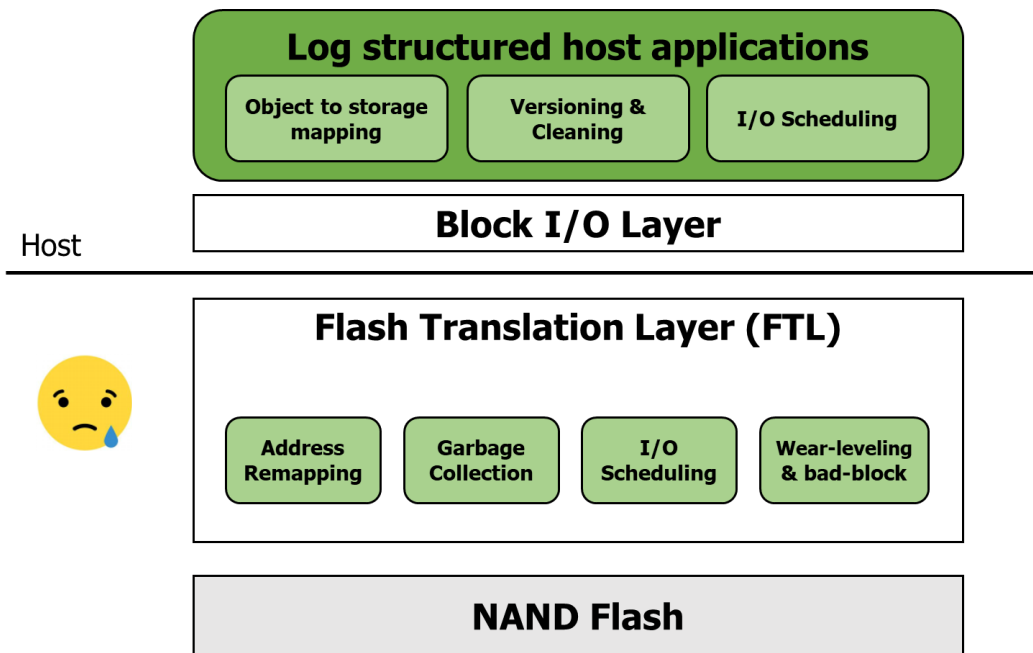
- Double logging between the FTL and the Host (double logging)
- Wastes hardware resource and increases write pressure to flash devices

Don't stack your Log on my Log

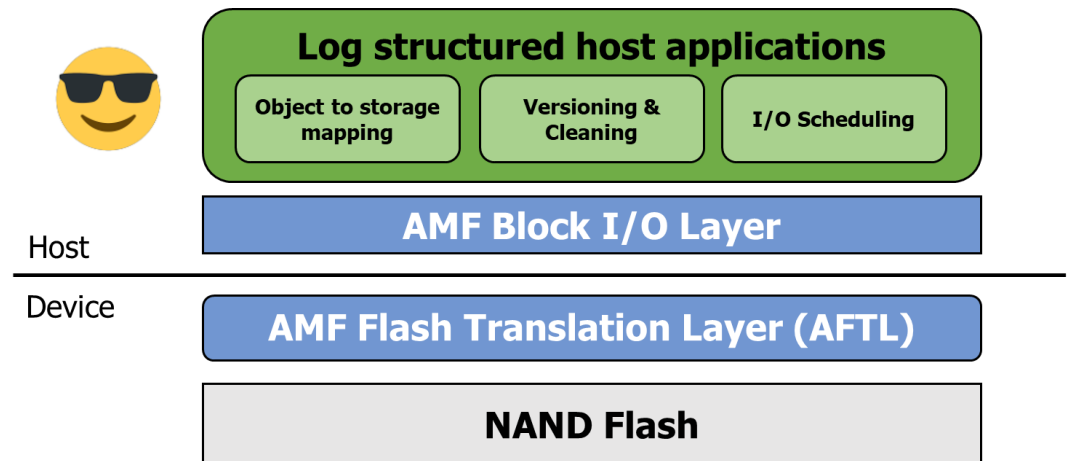
Jingpei Yang, Ned Plasson, Greg Gillis, Nisha Talagala, Swaminathan Sundararaman
SanDisk Corporation



Device-managed Flash



Application-managed Flash



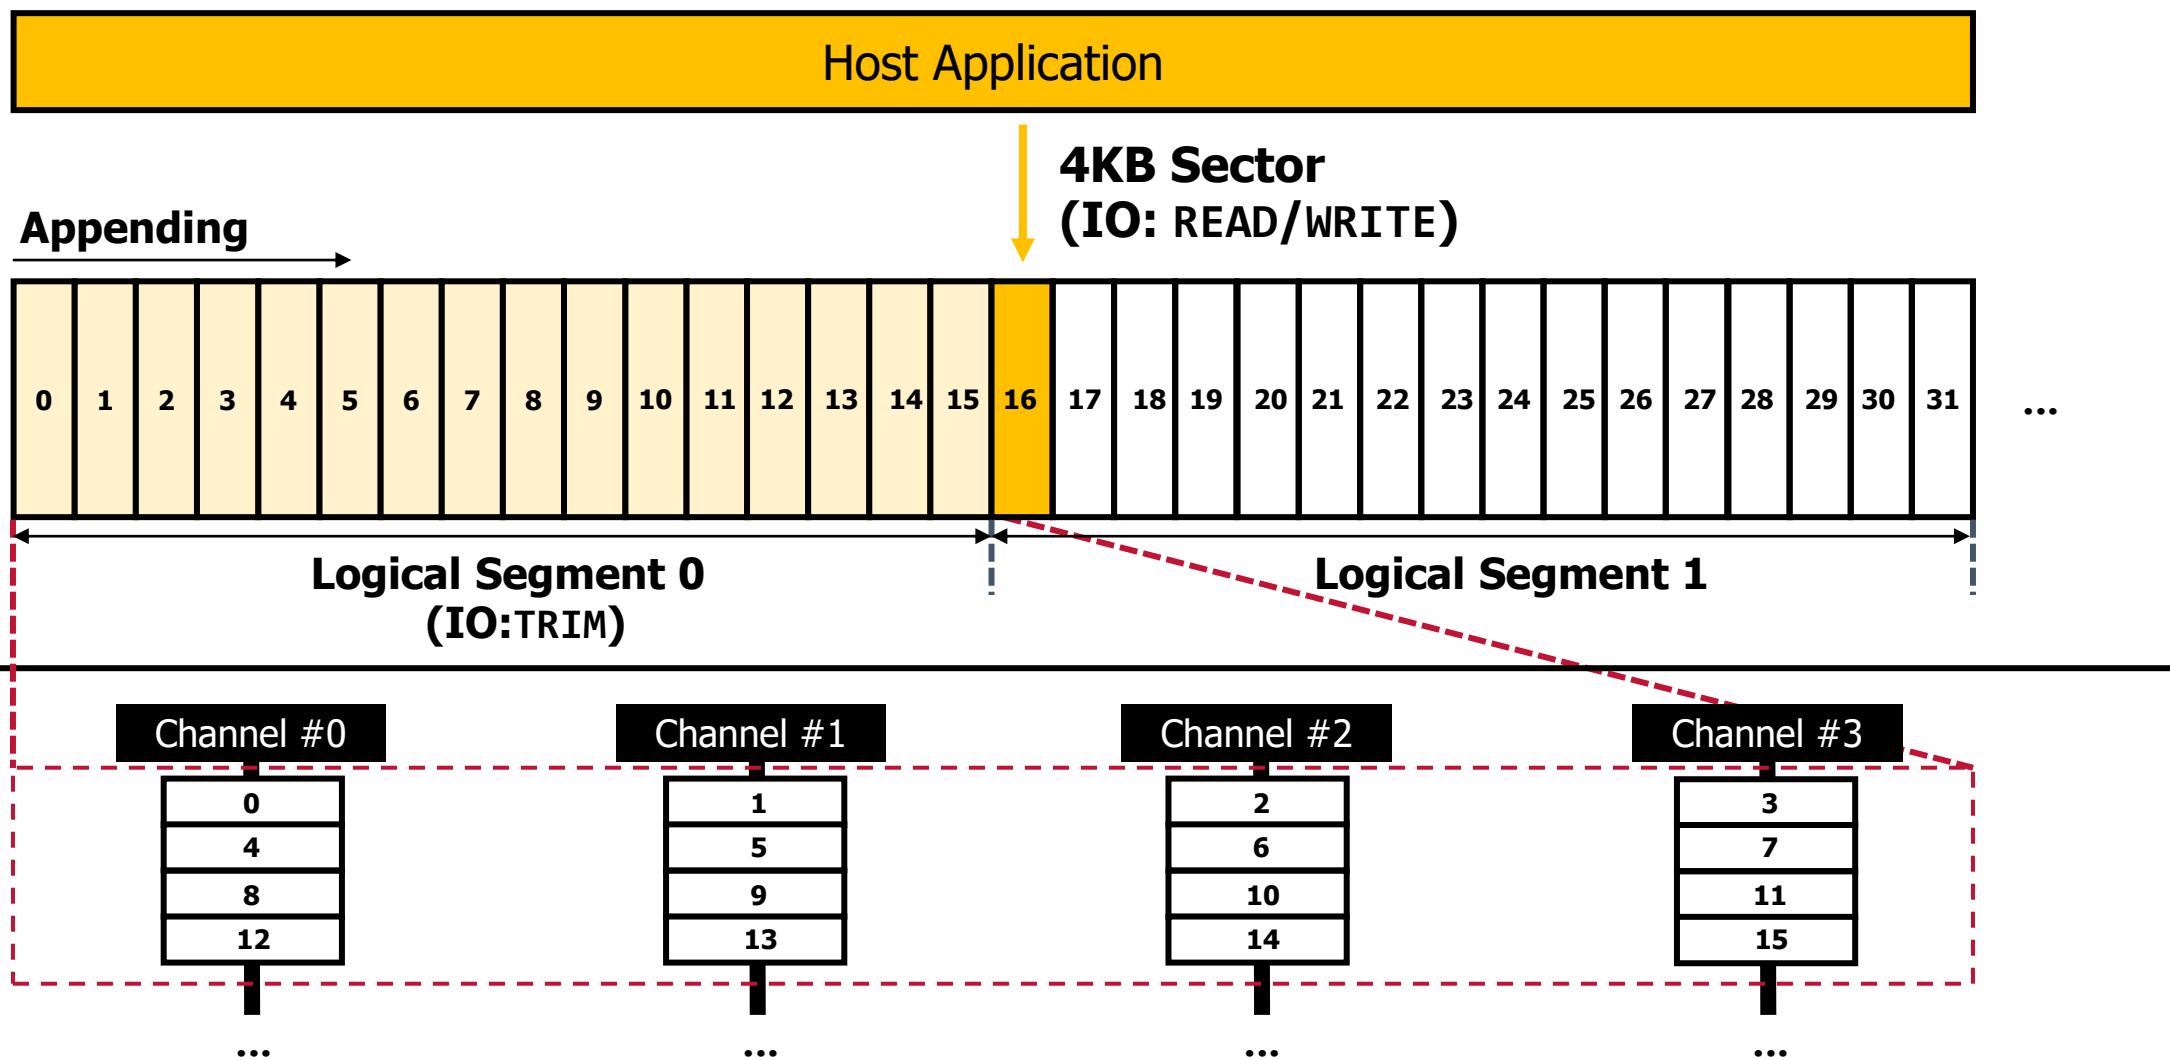
*moves flash management
from the device to applications*

Application-Managed Flash

- 1) Block I/O interface
- 2) ALFS (AMF Log-structured File System)
- 3) AFTL (AMF Flash Translation Layer)


Block I/O Interface

- AMF I/O is based on conventional block I/O



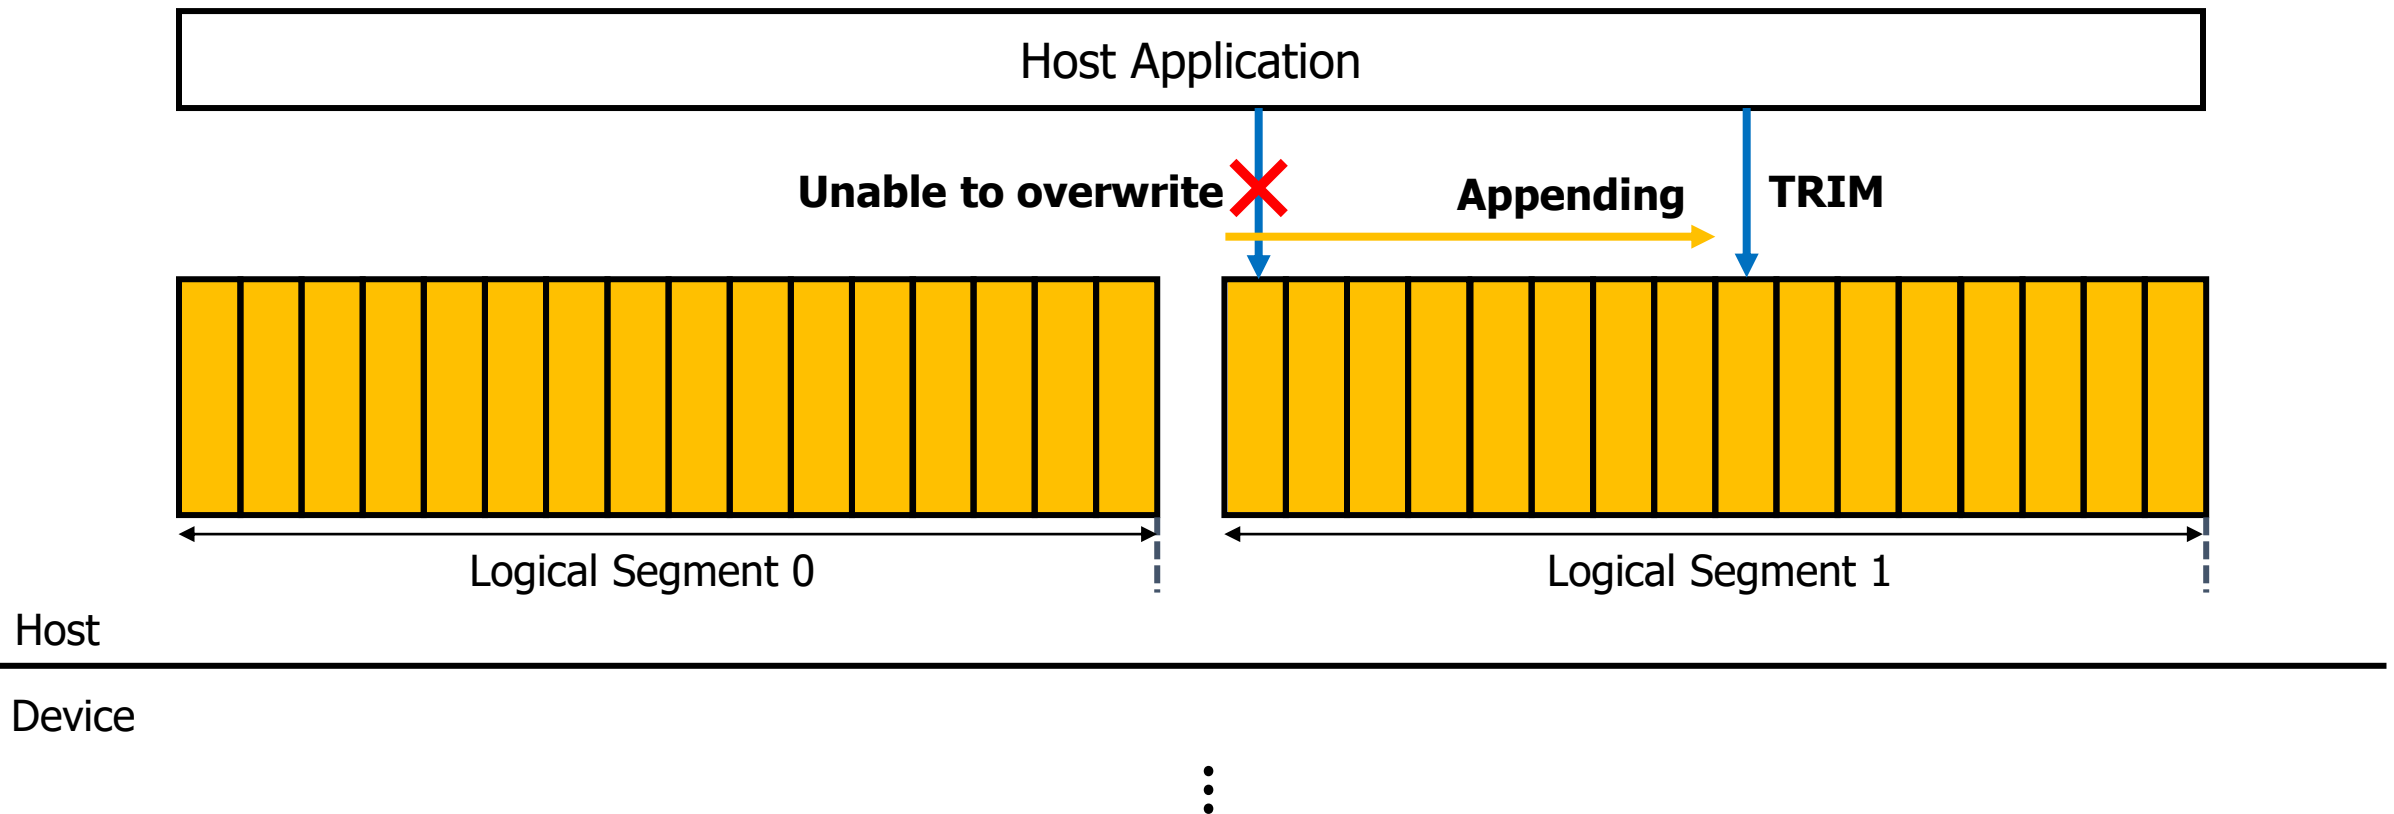
Block I/O Interface

- Append-only Write



Block I/O Interface

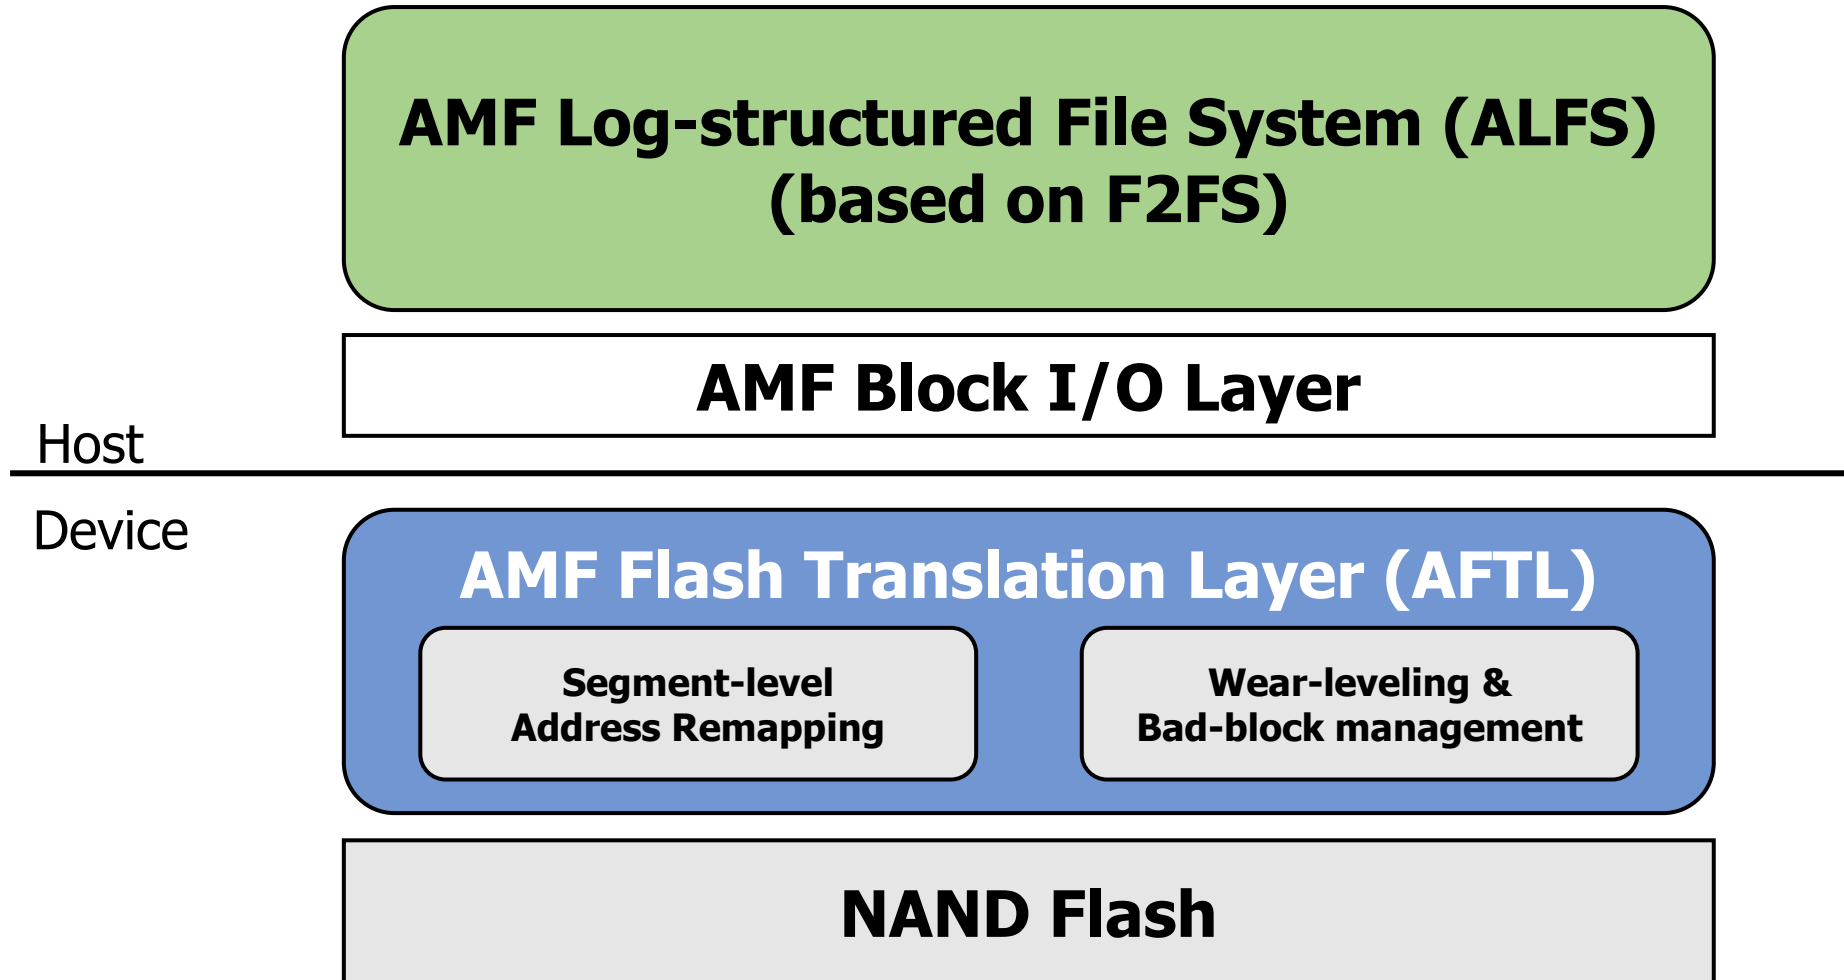
- Only sequential writes with no in-place updates
- Sectors can be reused after the segment has been deallocated by TRIM



Application-Managed Flash

- 1) Block I/O interface
- 2) **ALFS (AMF Log-structured File System)**
- 3) AFTL (AMF Flash Translation Layer)

ALFS (AMF Log-structured File System)



ALFS (AMF Log-structured File System)

- File Layout and Operation

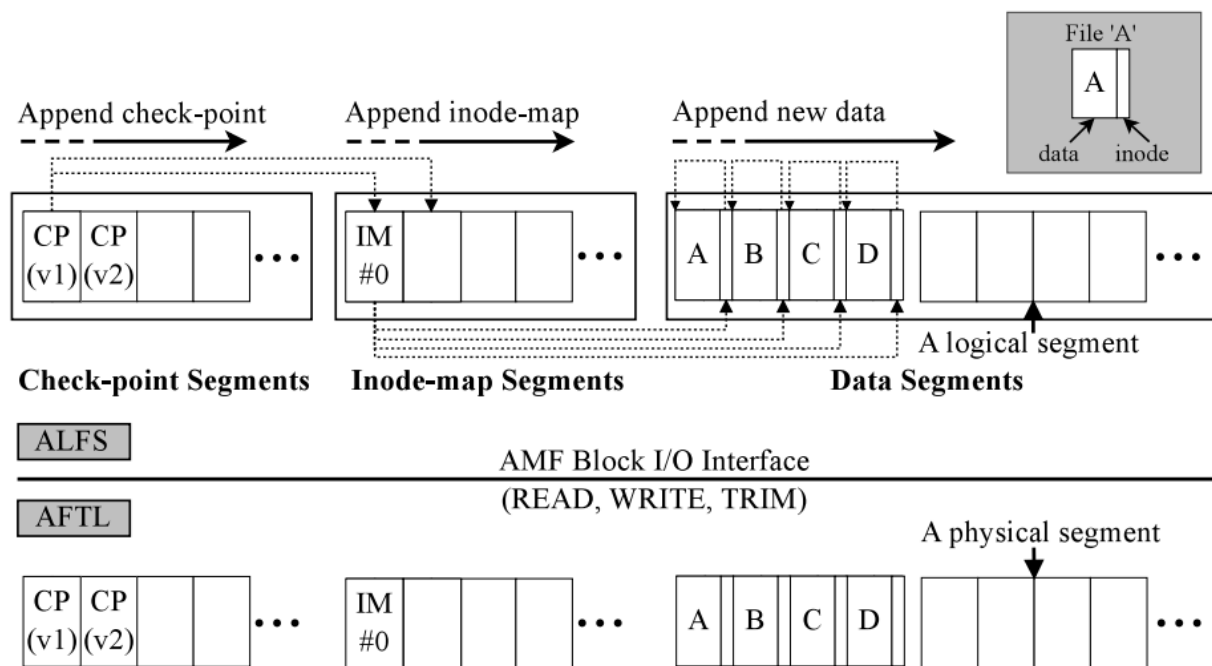
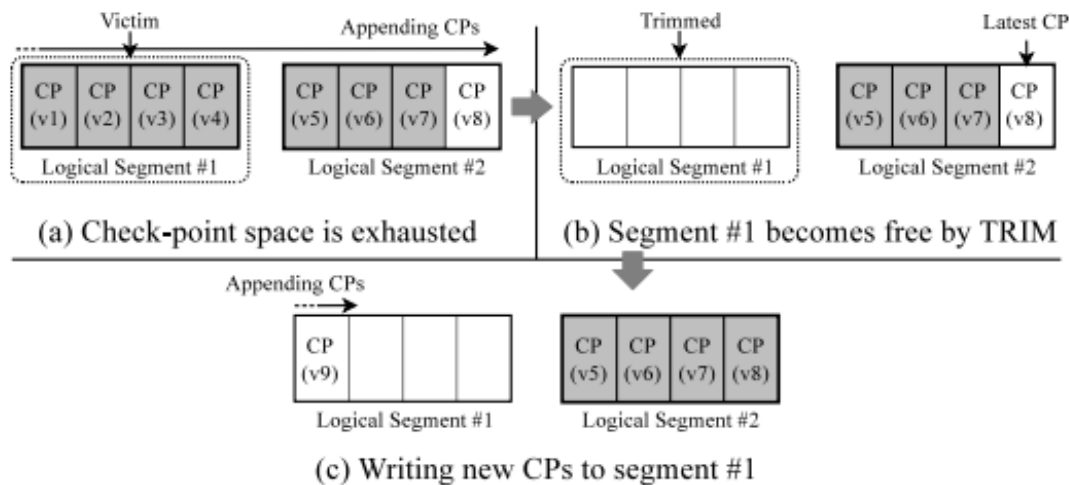


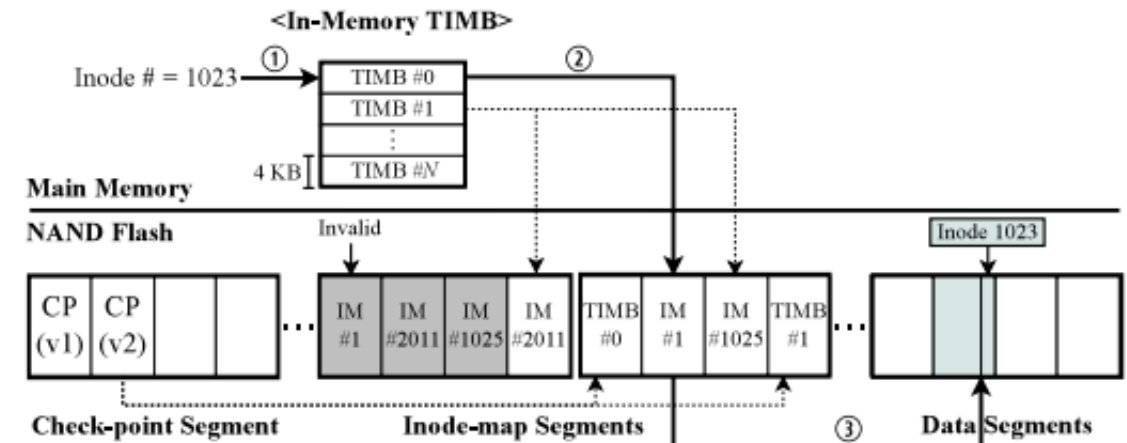
Figure 2: The upper figure illustrates the logical layout of ALFS. There is an initial check-point CP(v1). Four files are appended to data segments along with their inodes in the following order: A, B, C and D. Then, an inode map IM#0 is written which points to the locations of the inodes of the files. Finally, the check-point CP(v2) is written to check-point segments. The bottom figure shows the physical segments corresponding to the logical segments. The data layout of a logical segment perfectly aligns with its physical segment.

ALFS (AMF Log-structured File System)

- Check-point segments for quick mount and recovery
- Inode-map segments for fast searches of inodes



Check-point segment handling

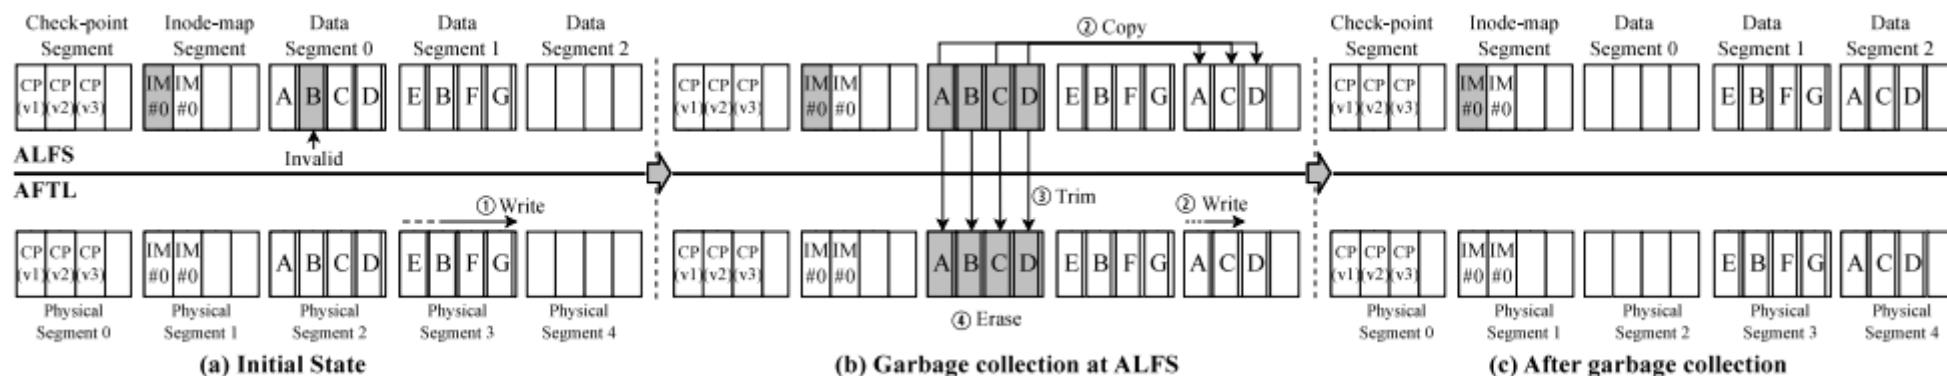


Managing inode-map block

ALFS (AMF Log-structured File System)

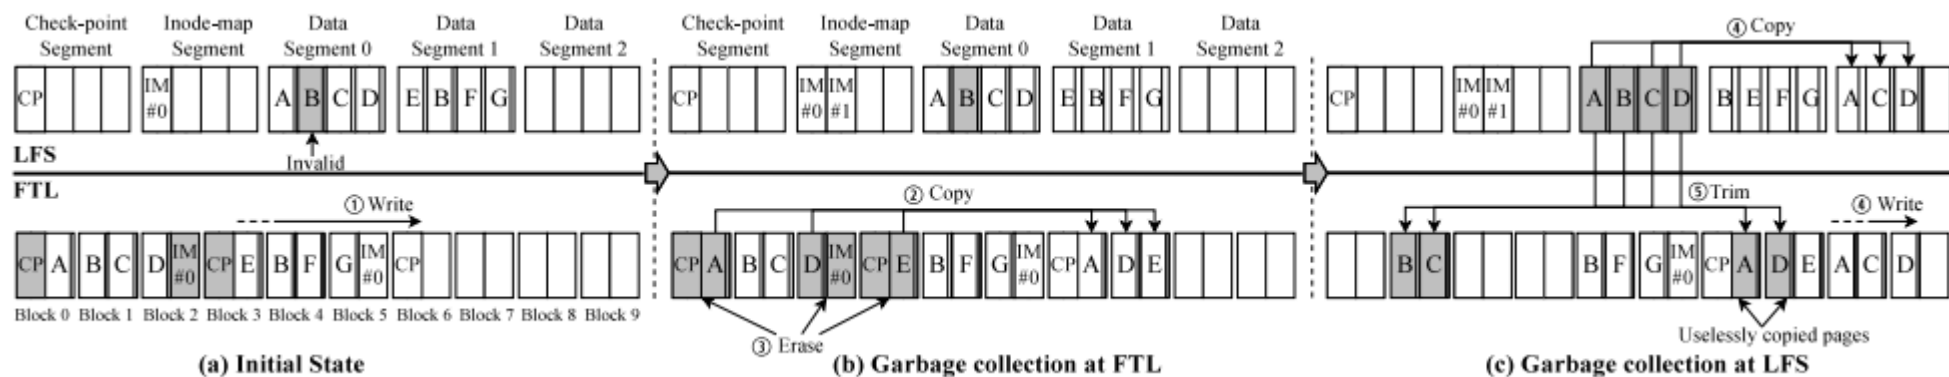
- Comparison with conventional LFS

ALFS
with
AFTL



3 page copies + 2 block erasures

LFS
with
FTL



6 page copies + 3 block erasures

Application-Managed Flash

- 1) Block I/O interface
- 2) ALFS (AMF Log-structured File System)
- 3) AFTL (AMF Flash Translation Layer)

AFTL (AMF flash Translation Layer)

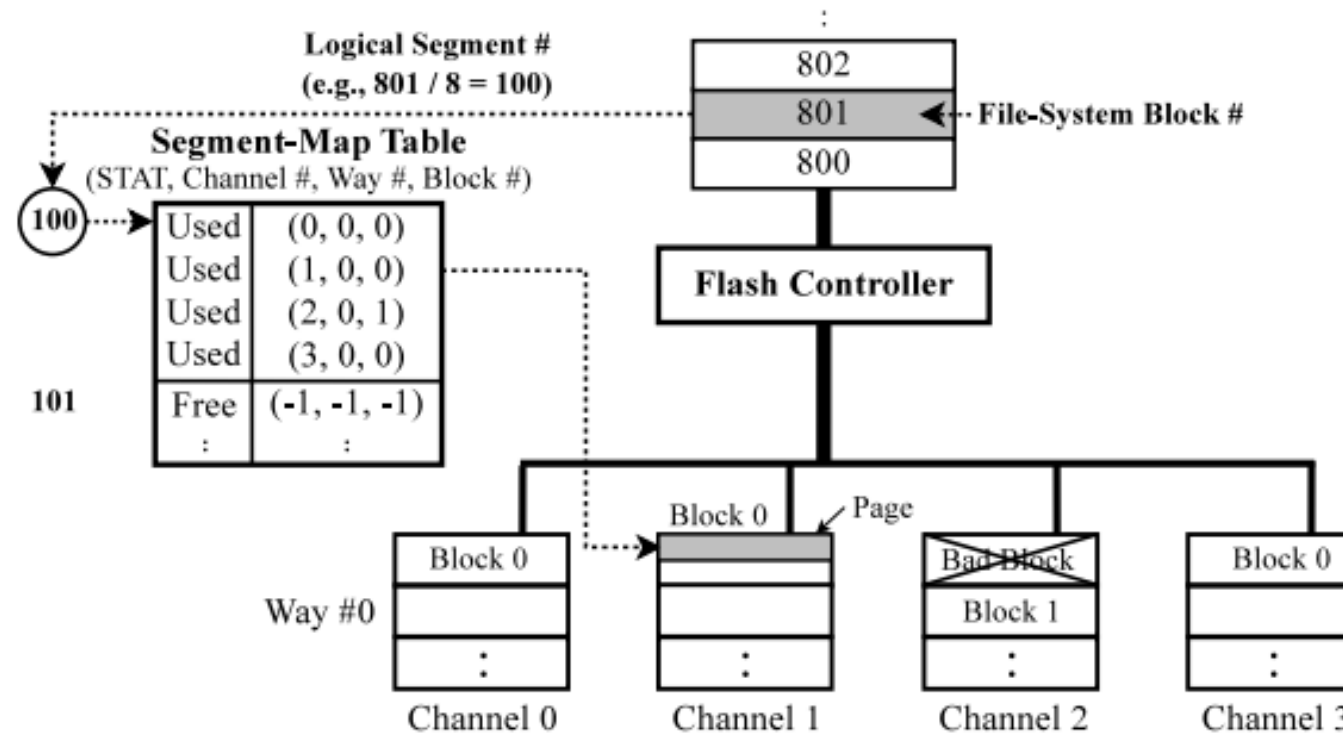


Figure 7: An example of how AFTL handles writes

Capacity	Block-level FTL	Hybrid FTL	Page-level FTL	AMF	
				AFTL	ALFS
512 GB	4 MB	96 MB	512 MB	4 MB	5.3 MB
1 TB	8 MB	186 MB	1 GB	8 MB	10.8 MB

Table 2: A summary of memory requirements

AFTL maintains a **smaller mapping table** than the page-level FTL

Data structures including TIMB need **tiny amount of host DRAM**

Performance measurment

- EXT4: only the FTL performed GC
- F2FS: both F2FS and the FTL performed GC
- AMF: only ALFS performed GC

- PFTL: no extra swapping I/Os
- AFTL: no extra swapping I/Os
- DFTL: incurred extra I/Os

Category	Workload	Description
File System	FIO	A synthetic I/O workload generator
	Postmark	A small and metadata intensive workload

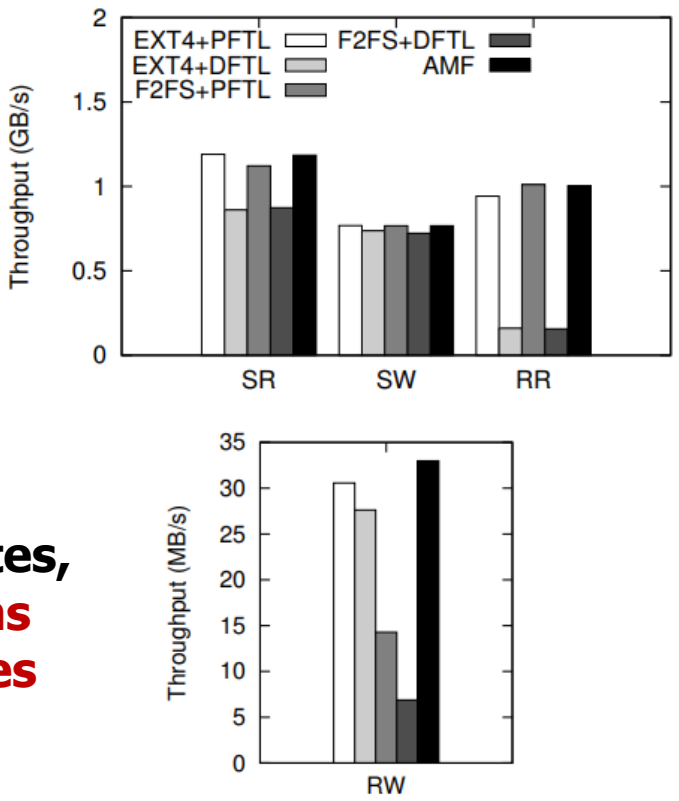


Figure 9: Experimental results with FIO

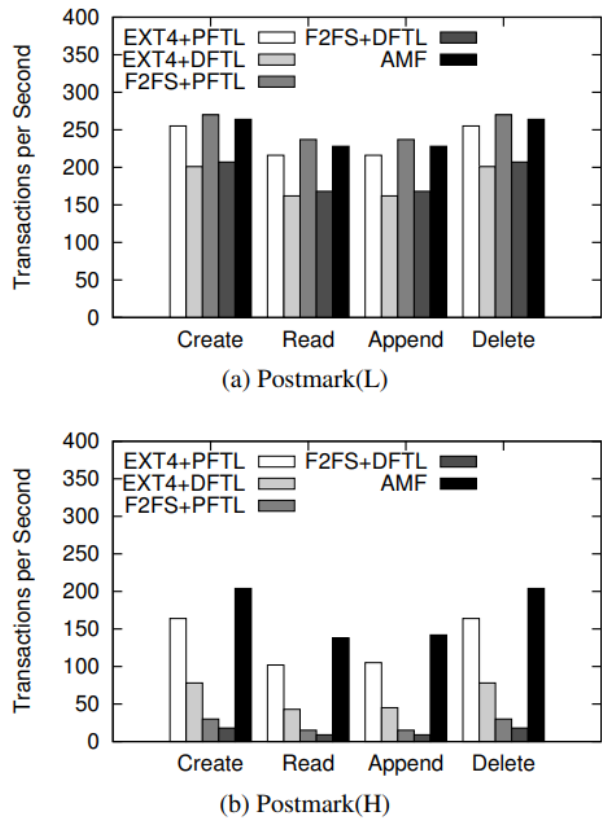


Figure 10: Experimental results with Postmark

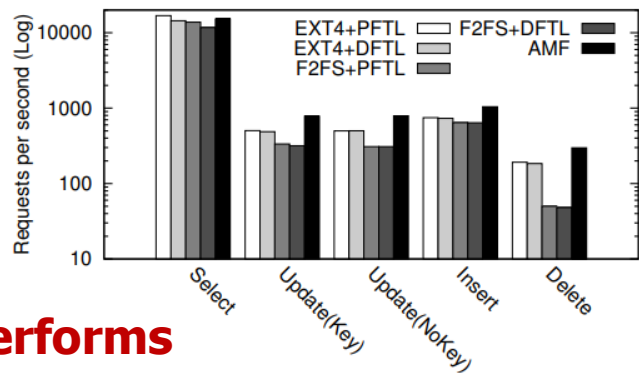
For random writes,
**AMF outperforms
all other schemes**

For the heavy workload,
**AMF achieves
the best performance**


Database	Non-Trans	A non-transactional DB workload
	OLTP	An OLTP workload
	TPC-C	A TPC-C workload

Hadoop	DFSIO	A HDFS I/O throughput test application
	TeraSort	A data sorting application
	WordCount	A word count application

AMF outperforms
all other schemes



(a) Non-Trans



(b) OLTP

(c) TPC-C

Figure 11: Experimental results with database apps.




Figure 12: Experimental results with Hadoop apps.

All five storage configurations
show **similar performance**




Figure 13: Erasure operations normalized to EXT4+PFTL

AMF incurs 28% fewer erase operations overall compared to F2FS+DFTL

**The CPU utilization of AMF
is similar to F2FS and EXT4**




Figure 15: CPU utilization (%)

**AMF has the shortest I/O
response times with small fluctuations**




Figure 16: Write latency (μsec)

AMF

- **New block I/O interface exposing flash storage as append-only segments**
- **Implementing new FTL scheme (AFTL) and a new file system (ALFS)**
- **Reducing DRAM in the flash controller by 128X, improving performance of the file system by 80%**

https://github.com/chamdoo/bdbm_drv.git

<https://github.com/sangwoojun/bluedbm.git>

Application-Managed Flash

Sungjin Lee, Ming Liu, Sangwoo Jun, and Shuotao Xu, Jihong Kim, Arvind
14th USENIX Conference on File and Storage Technologies (FAST '16). 2016.

Thank You!

2021. 06. 02

Presentation by Han, Yejin

hyj0225@dankook.ac.kr